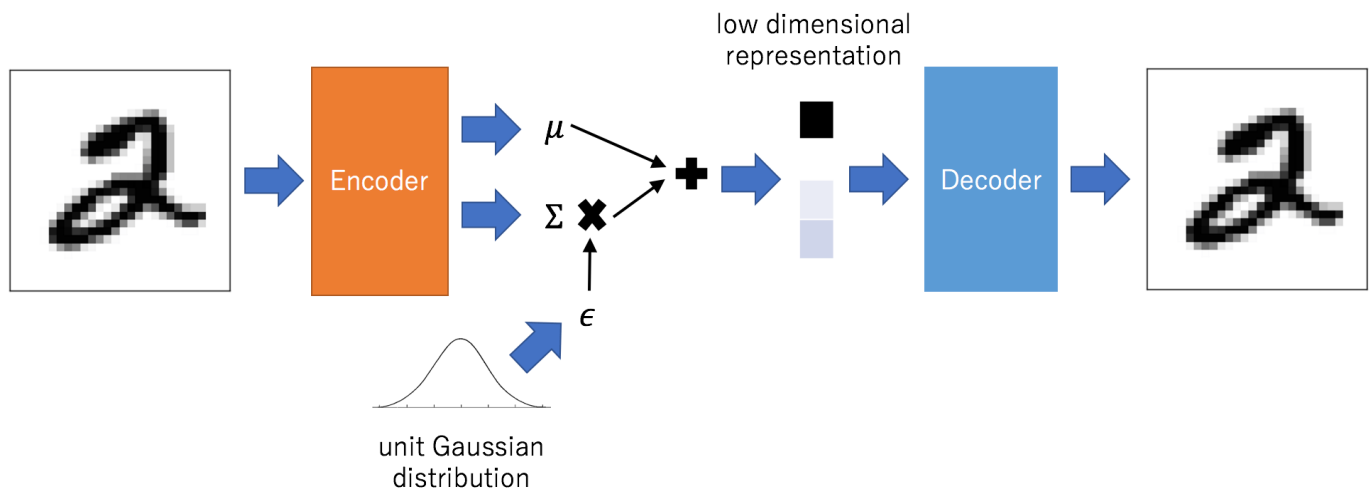


# Problem 1 - Variational Auto-Encoder (VAE)

Variational Auto-Encoders (VAEs) are a widely used class of generative models. They are simple to implement and, in contrast to other generative model classes like Generative Adversarial Networks (GANs, see Problem 2), they optimize an explicit maximum likelihood objective to train the model. Finally, their architecture makes them well-suited for unsupervised representation learning, i.e., learning low-dimensional representations of high-dimensional inputs, like images, with only self-supervised objectives (data reconstruction in the case of VAEs).



(image source: <https://mlexplained.com/2017/12/28/an-intuitive-explanation-of-variational-autoencoders-vaes-part-1>)

**By working on this problem you will learn and practice the following steps:**

1. Set up a data loading pipeline in PyTorch.
2. Implement, train and visualize an auto-encoder architecture.
3. Extend your implementation to a variational auto-encoder.
4. Learn how to tune the critical beta parameter of your VAE.
5. Inspect the learned representation of your VAE.
6. Extend VAE's generative capabilities by conditioning it on the label you wish to generate.

**Note:** For faster training of the models in this assignment you can enable GPU support in this Colab. Navigate to "Runtime" --> "Change Runtime Type" and set the "Hardware Accelerator" to "GPU". However, you might hit compute limits of the colab free edition. Hence, you might want to debug locally (e.g. in a jupyter notebook) or in a CPU-only runtime on colab.

## 1. MNIST Dataset

We will perform all experiments for this problem using the [MNIST dataset](#), a standard dataset of handwritten digits. The main benefits of this dataset are that it is small and relatively easy to model. It therefore allows for quick experimentation and serves as initial test bed in many papers.

Another benefit is that it is so widely used that PyTorch even provides functionality to automatically download it.

Let's start by downloading the data and visualizing some samples.

```
import matplotlib.pyplot as plt
%matplotlib inline

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

```
import torch
import torchvision
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu') # use GPU if
available
print(f"Using device: {device}")

# this will automatically download the MNIST training set
mnist_train = torchvision.datasets.MNIST(root='./data',
                                         train=True,
                                         download=True,
                                         transform=torchvision.transforms.ToTensor())
print("\n Download complete! Downloaded {} training
examples!".format(len(mnist_train)))
```

Using device: cuda:0

Download complete! Downloaded 60000 training examples!

```
/home/jingmin/miniconda3/envs/adaptor/lib/python3.8/site-
packages/torchvision/datasets/mnist.py:498: UserWarning: The given NumPy array is not
writeable, and PyTorch does not support non-writeable tensors. This means you can write
to the underlying (supposedly non-writeable) NumPy array using the tensor. You may want
to copy the array to protect its data or make it writeable before converting it to a
tensor. This type of warning will be suppressed for the rest of this program.
(triggered internally at /pytorch/torch/csrc/utils/tensor_numpy.cpp:180.)
    return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)
```

```
from numpy.random.mtrand import sample
import matplotlib.pyplot as plt
import numpy as np

# Let's display some of the training samples.
sample_images = []
randomize = False # set to False for debugging
num_samples = 5 # simple data sampling for now, later we will use proper DataLoader
```


```

if randomize:
    sample_idx = np.random.randint(low=0,high=len(mnist_train), size=num_samples)
else:
    sample_idx = list(range(num_samples))

for idx in sample_idx:
    sample = mnist_train[idx]
    # print(f"Tensor w/ shape {sample[0][0].detach().cpu().numpy().shape} and label {sample[1]}")
    sample_images.append(sample[0][0].data.cpu().numpy())
    # print(sample_images[0]) # Values are in [0, 1]

fig = plt.figure(figsize = (10, 50))
ax1 = plt.subplot(111)
ax1.imshow(np.concatenate(sample_images, axis=1), cmap='gray')
plt.show()

```



## 2. Auto-Encoder

Before implementing the full VAE, we will first implement an **auto-encoder architecture**. Auto-encoders feature the same encoder-decoder architecture as VAEs and therefore also learn a low-dimensional representation of the input data without supervision. In contrast to VAEs they are **fully deterministic** models and do not employ variational inference for optimization.

The **architecture** is very simple: we will encode the input image into a low-dimensional representation using fully connected layers for the encoder. This results in a low-dimensional representation of the input image. This representation will get decoded back into the dimensionality of the input image using a decoder network that mirrors the architecture of the encoder. The whole model is trained by **minimizing a reconstruction loss** between the input and the decoded image.

Intuitively, the **auto-encoder needs to compress the information contained in the input image** into a much lower dimensional representation (e.g.  $28 \times 28 = 784$ px vs.  $n_z$  embedding dimensions for our MNIST model). This is possible since the information captured in the pixels is *highly redundant*. E.g. encoding an MNIST image requires  $< 4$  bits to encode which of the 10 possible digits is displayed and a few additional bits to capture information about shape and orientation. This is much less than the  $255^{28 \cdot 28}$  bits of information that could be theoretically captured in the input image.

Learning such a **compressed representation can make downstream task learning easier**. For example, learning to add two numbers based on the inferred digits is much easier than performing the task based on two piles of pixel values that depict the digits.

In the following, we will first define the architecture of encoder and decoder and then train the auto-encoder model.

## Defining the Auto-Encoder Architecture [6pt]

```
import torch.nn as nn

# Probl-1: Let's define encoder and decoder networks

class Encoder(nn.Module):
    def __init__(self, nz, input_size):
        super().__init__()
        self.input_size = input_size
        ##### TODO #####
        # Create the network architecture using a nn.Sequential module wrapper. #
        # Encoder Architecture: #
        # - input_size -> 256 #
        # - ReLU #
        # - 256 -> 64 #
        # - ReLU #
        # - 64 -> nz #
        # HINT: Verify the shapes of intermediate layers by running partial networks #
        # (with the next notebook cell) and visualizing the output shapes. #
        #####
        self.net = nn.Sequential(
            nn.Linear(input_size, 256),
            nn.ReLU(),
            nn.Linear(256, 64),
            nn.ReLU(),
            nn.Linear(64, nz)
        )
        ##### END TODO #####

    def forward(self, x):
        return self.net(x)

class Decoder(nn.Module):
    def __init__(self, nz, output_size):
        super().__init__()
        self.output_size = output_size
        ##### TODO #####
        # Create the network architecture using a nn.Sequential module wrapper. #
        # Decoder Architecture (mirrors encoder architecture): #
        # - nz -> 64 #
        # - ReLU #
```

```

# - 64 -> 256 #
# - ReLU #
# - 256 -> output_size #
#####
self.net = nn.Sequential(
    nn.Linear(nz, 64),
    nn.ReLU(),
    nn.Linear(64, 256),
    nn.ReLU(),
    nn.Linear(256, output_size),
    nn.Sigmoid()
)
##### END TODO #####

def forward(self, z):
    return self.net(z).reshape(-1, 1, self.output_size)

```

## Testing the Auto-Encoder Forward Pass

```

# To test your encoder/decoder, let's encode/decode some sample images
# first, make a PyTorch DataLoader object to sample data batches
batch_size = 64
nworkers = 2 # number of workers used for efficient data loading

#####
# Create a PyTorch DataLoader object for efficiently generating training batches. #
# Make sure that the data loader automatically shuffles the training dataset. #
# Consider only *full* batches of data, to avoid torch errors. #
# The DataLoader wraps the MNIST dataset class we created earlier. #
# Use the given batch_size and number of data loading workers when creating #
# the DataLoader. https://pytorch.org/docs/stable/data.html #
#####
mnist_data_loader = torch.utils.data.DataLoader(mnist_train,
                                                batch_size=batch_size,
                                                shuffle=True,
                                                num_workers=nworkers,
                                                drop_last=True)
#####

# now we can run a forward pass for encoder and decoder and check the produced shapes
in_size = out_size = 28*28 # image size
nz = 32 # dimensionality of the learned embedding
encoder = Encoder(nz=nz, input_size=in_size)
decoder = Decoder(nz=nz, output_size=out_size)
for sample_img, sample_label in mnist_data_loader: # loads a batch of data
    input = sample_img.reshape([batch_size, in_size])
    print(f'{sample_img.shape=}, {type(sample_img)}, {input.shape=}')
    enc = encoder(input)

```

```

print(f"Shape of encoding vector (should be [batch_size, nz]): {enc.shape}")
dec = decoder(enc)
print("Shape of decoded image (should be [batch_size, 1, out_size]):
{}".format(dec.shape))
break

del input, enc, dec, encoder, decoder, nworkers # remove to avoid confusion later

```

```

sample_img.shape=torch.Size([64, 1, 28, 28]), <class 'torch.Tensor'>,
input.shape=torch.Size([64, 784])
Shape of encoding vector (should be [batch_size, nz]): torch.Size([64, 32])
Shape of decoded image (should be [batch_size, 1, out_size]): torch.Size([64, 1, 784]).

```

Now that we defined encoder and decoder network our architecture is nearly complete. However, before we start training, we can wrap encoder and decoder into an auto-encoder class for easier handling.

```

class AutoEncoder(nn.Module):
    def __init__(self, nz):
        super().__init__()
        self.encoder = Encoder(nz=nz, input_size=in_size)
        self.decoder = Decoder(nz=nz, output_size=out_size)

    def forward(self, x):
        enc = self.encoder(x)
        return self.decoder(enc)

    def reconstruct(self, x):
        """Only used later for visualization."""
        enc = self.encoder(x)
        flattened = self.decoder(enc)
        image = flattened.reshape(-1, 28, 28)
        return image

```

## Setting up the Auto-Encoder Training Loop [6pt]

After implementing the network architecture, we can now set up the training loop and run training.

```

# Probl-2
epochs = 10
learning_rate = 1e-3

# build AE model
print(f'Device available {device}')
ae_model = AutoEncoder(nz).to(device) # transfer model to GPU if available
ae_model = ae_model.train() # set model in train mode (eg batchnorm params get
updated)

```

```

# build optimizer and loss function
##### TODO #####
# Build the optimizer and loss classes. For the loss you can use a loss layer #
# from the torch.nn package. We recommend binary cross entropy. #
# HINT: We will use the Adam optimizer (learning rate given above, otherwise #
#         default parameters). #
# NOTE: We could also use alternative losses like MSE and cross entropy, depending #
#         on the assumptions we are making about the output distribution. #
#####
optimizer = torch.optim.Adam(ae_model.parameters(), learning_rate)
loss_func = nn.BCELoss(reduction='mean')
# loss_func = nn.MSELoss()
##### END TODO #####

train_it = 0
for ep in range(epochs):
    print("Run Epoch {}".format(ep))
    ##### TODO #####
    #####
    # Implement the main training loop for the auto-encoder model.
    #
    # HINT: Your training loop should sample batches from the data loader, run the
    #
    #         forward pass of the AE, compute the loss, perform the backward pass and
    #
    #         perform one gradient step with the optimizer.
    #
    # HINT: Don't forget to erase old gradients before performing the backward pass.
    #

    #####
    for sample_img, _ in mnist_data_loader: # _ refers to y_label
        sample_img = sample_img.to(device)
        sample_img = sample_img.reshape([batch_size, in_size])
        optimizer.zero_grad()
        # falten X_train: [64, 28, 28] -> [64, 784]
        y_pred = ae_model(sample_img)
        # print(y_pred.shape, X_train.unsqueeze(1).shape)
        # loss between reconstructed img and original img
        # rec_loss = loss_func(y_pred, X_train.unsqueeze(1))
        rec_loss = loss_func(y_pred.view(batch_size, -1), sample_img) # unsqueeze
X_train: [64, 784] -> [64, 1, 784]
        rec_loss.backward()
        optimizer.step()
        if train_it % 100 == 0:
            print("It {}: Reconstruction Loss: {}".format(train_it, rec_loss))
            train_it += 1
    ##### END TODO #####

```

```
print("Done!")
del epochs, learning_rate, sample_img, train_it, rec_loss #, opt
```

Device available cuda:0

Run Epoch 0

It 0: Reconstruction Loss: 0.6941973567008972  
It 100: Reconstruction Loss: 0.2600642442703247  
It 200: Reconstruction Loss: 0.21377862989902496  
It 300: Reconstruction Loss: 0.19295649230480194  
It 400: Reconstruction Loss: 0.1773861050605774  
It 500: Reconstruction Loss: 0.15200963616371155  
It 600: Reconstruction Loss: 0.15025803446769714  
It 700: Reconstruction Loss: 0.1468232423067093  
It 800: Reconstruction Loss: 0.15159788727760315  
It 900: Reconstruction Loss: 0.12887603044509888

Run Epoch 1

It 1000: Reconstruction Loss: 0.13041232526302338  
It 1100: Reconstruction Loss: 0.13362430036067963  
It 1200: Reconstruction Loss: 0.12781736254692078  
It 1300: Reconstruction Loss: 0.12402553856372833  
It 1400: Reconstruction Loss: 0.12981130182743073  
It 1500: Reconstruction Loss: 0.11978445202112198  
It 1600: Reconstruction Loss: 0.11579222232103348  
It 1700: Reconstruction Loss: 0.11232109367847443  
It 1800: Reconstruction Loss: 0.10930512845516205

Run Epoch 2

It 1900: Reconstruction Loss: 0.12100382894277573  
It 2000: Reconstruction Loss: 0.11269942671060562  
It 2100: Reconstruction Loss: 0.11609923839569092  
It 2200: Reconstruction Loss: 0.11043394356966019  
It 2300: Reconstruction Loss: 0.1102084144949913  
It 2400: Reconstruction Loss: 0.1063244491815567  
It 2500: Reconstruction Loss: 0.10522106289863586  
It 2600: Reconstruction Loss: 0.10785263031721115  
It 2700: Reconstruction Loss: 0.1129796952009201  
It 2800: Reconstruction Loss: 0.11421039700508118

Run Epoch 3

It 2900: Reconstruction Loss: 0.10506419837474823  
It 3000: Reconstruction Loss: 0.10360829532146454  
It 3100: Reconstruction Loss: 0.10624536126852036  
It 3200: Reconstruction Loss: 0.10837209969758987  
It 3300: Reconstruction Loss: 0.10888055711984634  
It 3400: Reconstruction Loss: 0.1086772009730339  
It 3500: Reconstruction Loss: 0.10291092842817307  
It 3600: Reconstruction Loss: 0.09583210200071335  
It 3700: Reconstruction Loss: 0.10236454010009766

Run Epoch 4

It 3800: Reconstruction Loss: 0.09569605439901352  
It 3900: Reconstruction Loss: 0.10292566567659378



It 4000: Reconstruction Loss: 0.09966414421796799  
It 4100: Reconstruction Loss: 0.10311181098222733  
It 4200: Reconstruction Loss: 0.10047150403261185  
It 4300: Reconstruction Loss: 0.10418491810560226  
It 4400: Reconstruction Loss: 0.09745653718709946  
It 4500: Reconstruction Loss: 0.09582802653312683  
It 4600: Reconstruction Loss: 0.0952438935637474

Run Epoch 5

It 4700: Reconstruction Loss: 0.09557317942380905  
It 4800: Reconstruction Loss: 0.09920232743024826  
It 4900: Reconstruction Loss: 0.09611573070287704  
It 5000: Reconstruction Loss: 0.10116905719041824  
It 5100: Reconstruction Loss: 0.09177839010953903  
It 5200: Reconstruction Loss: 0.0999147817492485  
It 5300: Reconstruction Loss: 0.09217223525047302  
It 5400: Reconstruction Loss: 0.09530997276306152  
It 5500: Reconstruction Loss: 0.0955212265253067  
It 5600: Reconstruction Loss: 0.09844547510147095

Run Epoch 6

It 5700: Reconstruction Loss: 0.08832038193941116  
It 5800: Reconstruction Loss: 0.09406760334968567  
It 5900: Reconstruction Loss: 0.0942428931593895  
It 6000: Reconstruction Loss: 0.09220677614212036  
It 6100: Reconstruction Loss: 0.09558732062578201  
It 6200: Reconstruction Loss: 0.08982578665018082  
It 6300: Reconstruction Loss: 0.09608685225248337  
It 6400: Reconstruction Loss: 0.09362149238586426  
It 6500: Reconstruction Loss: 0.09448808431625366

Run Epoch 7

It 6600: Reconstruction Loss: 0.09203379601240158  
It 6700: Reconstruction Loss: 0.09271685779094696  
It 6800: Reconstruction Loss: 0.08378888666629791  
It 6900: Reconstruction Loss: 0.08798205107450485  
It 7000: Reconstruction Loss: 0.08749934285879135  
It 7100: Reconstruction Loss: 0.08796916157007217  
It 7200: Reconstruction Loss: 0.09372701495885849  
It 7300: Reconstruction Loss: 0.0945957824587822  
It 7400: Reconstruction Loss: 0.0908859372138977

Run Epoch 8

It 7500: Reconstruction Loss: 0.08930474519729614  
It 7600: Reconstruction Loss: 0.09202592819929123  
It 7700: Reconstruction Loss: 0.0878760814666748  
It 7800: Reconstruction Loss: 0.08705627173185349  
It 7900: Reconstruction Loss: 0.09490896761417389  
It 8000: Reconstruction Loss: 0.08935634046792984  
It 8100: Reconstruction Loss: 0.09234100580215454  
It 8200: Reconstruction Loss: 0.09026658535003662  
It 8300: Reconstruction Loss: 0.0871456116437912  
It 8400: Reconstruction Loss: 0.09165652841329575

```
Run Epoch 9
It 8500: Reconstruction Loss: 0.08753928542137146
It 8600: Reconstruction Loss: 0.09677305072546005
It 8700: Reconstruction Loss: 0.08558925241231918
It 8800: Reconstruction Loss: 0.08861494809389114
It 8900: Reconstruction Loss: 0.0951785147190094
It 9000: Reconstruction Loss: 0.083978570997715
It 9100: Reconstruction Loss: 0.08792658895254135
It 9200: Reconstruction Loss: 0.09171804040670395
It 9300: Reconstruction Loss: 0.08485647290945053
Done!
```

## Verifying reconstructions

Now that we trained the auto-encoder we can visualize some of the reconstructions on the test set to verify that it is converged and did not overfit. **Before continuing, make sure that your auto-encoder is able to reconstruct these samples near-perfectly.**

```
# visualize test data reconstructions
def vis_reconstruction(model, randomize=False):
    # download MNIST test set + build Dataset object
    mnist_test = torchvision.datasets.MNIST(root='./data',
                                             train=False,
                                             download=True,

transform=torchvision.transforms.ToTensor())
    model.eval()      # set model in evaluation mode (eg freeze batchnorm params)
    num_samples = 5
    if randomize:
        sample_idx = np.random.randint(low=0, high=len(mnist_test), size=num_samples)
    else:
        sample_idx = list(range(num_samples))


    input_imgs, test_reconstructions = [], []
    for idx in sample_idx:
        sample = mnist_test[idx]
        input_img = np.asarray(sample[0])
        input_flat = input_img.reshape(784)
        reconstruction = model.reconstruct(torch.tensor(input_flat, device=device))

        input_imgs.append(input_img[0])
        test_reconstructions.append(reconstruction[0].data.cpu().numpy())
        # print(f'{input_img[0].shape}\t{reconstruction.shape}')

    fig = plt.figure(figsize = (20, 50))
    ax1 = plt.subplot(111)
    ax1.imshow(np.concatenate([np.concatenate(input_imgs, axis=1),
                               np.concatenate(test_reconstructions, axis=1)], axis=0),
               cmap='gray')
```

```
plt.show()
```

```
vis_reconstruction(ae_model, randomize=False) # set randomize to False for debugging
```



## Sampling from the Auto-Encoder [2pt]

To test whether the auto-encoder is useful as a generative model, we can use it like any other generative model: draw embedding samples from a prior distribution and decode them through the decoder network. We will choose a unit Gaussian prior to allow for easy comparison to the VAE later.

```
# we will sample N embeddings, then decode and visualize them
def vis_samples(model):
    ##### TODO
    #####
    # Probl-3 Sample embeddings from a diagonal unit Gaussian distribution and decode
    them #
    # using the model.
    #
    # HINT: The sampled embeddings should have shape [batch_size, nz]. Diagonal unit
    #
    # Gaussians have mean 0 and a covariance matrix with ones on the diagonal
    #
    # and zeros everywhere else.
    #
    # HINT: If you are unsure whether you sampled the correct distribution, you can
    #
    # sample a large batch and compute the empirical mean and variance using the
    #
    # .mean() and .var() functions.
    #
```

```


# HINT: You can directly use model.decoder() to decode the samples.
#

#####
sampled_embeddings = torch.randn((batch_size, nz)).to(device) # sample batch of
embedding from prior
# print(sampled_embeddings.mean(), sampled_embeddings.var())
# decoder output images for sampled embeddings
decoded_samples = model.decoder(sampled_embeddings).reshape(-1, 1, 28, 28)
# print(decoded_samples.shape)
##### END TODO
#####

fig = plt.figure(figsize = (10, 10))
ax1 = plt.subplot(111)
ax1.imshow(torchvision.utils.make_grid(decoded_samples[:16], nrow=4, pad_value=1.)\
          .data.cpu().numpy().transpose(1, 2, 0), cmap='gray')
plt.show()

vis_samples(ae_model)

```



**Prob1-3 continued: Inline Question: Describe your observations, why do you think they occur?**  
**[2pt]** (max 150 words)

**Answer:** The encoded latent space may not be regular, which can make it difficult for AE model to perform smooth and meaningful interpolations between different points in the space. This irregularity and discontinuity can result from various factors such as the distribution of data, model architecture, and the dimension of the latent space. Furthermore, AE cannot learn the underlying distribution of the input data, which can limit their ability to generate high-quality samples from the latent space. Thus, the results decoded from Gaussain distribution are low-quality.

### 3. Variational Auto-Encoder (VAE)

Variational auto-encoders use a very similar architecture to deterministic auto-encoders, but are inherently stochastic models, i.e. we perform a stochastic sampling operation during the forward pass, leading to different different outputs every time we run the network for the same input. This sampling is required to optimize the VAE objective also known as the evidence lower bound (ELBO):

$$p(x) > \underbrace{\mathbb{E}_{z \sim q(z|x)} p(x|z)}_{\text{reconstruction}} - \underbrace{D_{\text{KL}}(q(z|x), p(z))}_{\text{prior divergence}}$$

Here,  $D_{\text{KL}}(q, p)$  denotes the Kullback-Leibler (KL) divergence between the posterior distribution  $q(z|x)$ , i.e. the output of our encoder, and  $p(z)$ , the prior over the embedding variable  $z$ , which we can choose freely.

For simplicity, we will choose a unit Gaussian prior again. The first term is the reconstruction term we already know from training the auto-encoder. When assuming a Gaussian output distribution for both encoder  $q(z|x)$  and decoder  $p(x|z)$  the objective reduces to:

$$\mathcal{L}_{\text{VAE}} = \sum_{x \sim \mathcal{D}} \mathcal{L}_{\text{rec}}(x, \hat{x}) - \beta \cdot D_{\text{KL}}(\mathcal{N}(\mu_q, \sigma_q), \mathcal{N}(0, I))$$

Here,  $\hat{x}$  is the reconstruction output of the decoder. In comparison to the auto-encoder objective, the VAE adds a regularizing term between the output of the encoder and a chosen prior distribution, effectively forcing the encoder output to not stray too far from the prior during training. As a result the decoder gets trained with samples that look pretty similar to samples from the prior, which will hopefully allow us to generate better images when using the VAE as a generative model and actually feeding it samples from the prior (as we have done for the AE before).

The coefficient  $\beta$  is a scalar weighting factor that trades off between reconstruction and regularization objective. We will investigate the influence of this factor in our experiments below.

If you need a refresher on VAEs you can check out this tutorial paper: <https://arxiv.org/abs/1606.05908>

## Reparametrization Trick

The sampling procedure inside the VAE's forward pass for obtaining a sample  $z$  from the posterior distribution  $q(z|x)$ , when implemented naively, is non-differentiable. However, since  $q(z|x)$  is parametrized with a Gaussian function, there is a simple trick to obtain a differentiable sampling operator, known as the *reparametrization trick*.

Instead of directly sampling  $z \sim \mathcal{N}(\mu_q, \sigma_q)$  we can "separate" the network's predictions and the random sampling by computing the sample as:

$$z = \mu_q + \sigma_q * \epsilon, \quad \epsilon \sim \mathcal{N}(0, I)$$

Note that in this equation, the sample  $z$  is computed as a deterministic function of the network's predictions  $\mu_q$  and  $\sigma_q$  and therefore allows to propagate gradients through the sampling procedure.

**Note:** While in the equations above the encoder network parametrizes the standard deviation  $\sigma_q$  of the Gaussian posterior distribution, in practice we usually parametrize the **logarithm of the standard deviation**  $\log \sigma_q$  for numerical stability. Before sampling  $z$  we will then exponentiate the network's output to obtain  $\sigma_q$ .

# Defining the VAE Model [7pt]

```
def kl_divergence(mu1, log_sigma1, mu2, log_sigma2):
    """Computes KL[p||q] between two Gaussians defined by [mu, log_sigma]."""
    return (log_sigma2 - log_sigma1) + (torch.exp(log_sigma1) ** 2 + (mu1 - mu2) ** 2)
    \
        / (2 * torch.exp(log_sigma2) ** 2) - 0.5

# Probl-4
class VAE(nn.Module):
    def __init__(self, nz, beta=1.0):
        super().__init__()
        self.beta = beta          # factor trading off between two loss components
        ##### TODO
        #####
        # Instantiate Encoder and Decoder.
        #
        # HINT: Remember that the encoder is now parametrizing a Gaussian
distribution's    #
        #         mean and log_sigma, so the dimensionality of the output needs to
        #
        #         double. The decoder works with an embedding sampled from this output.
        #
        #####
        self.nz = nz
        self.encoder = Encoder(2 * self.nz, in_size)
        self.decoder = Decoder(self.nz, out_size)
        ##### END TODO
        #####

    def forward(self, x):
        ##### TODO
        #####
        # Implement the forward pass of the VAE.
        #
        # HINT: Your code should implement the following steps:
        #
        #         1. encode input x, split encoding into mean and log_sigma of
Gaussian    #
        #         2. sample z from inferred posterior distribution using
        #
        #         reparametrization trick
        #
        #         3. decode the sampled z to obtain the reconstructed image
        #
        #####
        q = self.encoder(x)    # encode x
```

```

mu = q[:, :self.nz] # split encoding
log_sigma = q[:, self.nz:]
# mu, log_sigma = torch.chunk(q, 2, dim=-1)
# reparametrization: z = mu + sigma * epsilon
eps = torch.randn(log_sigma.shape).to(device)
# eps = torch.normal(mean=torch.zeros_like(mu), std=torch.ones_like(log_sigma))
# sample latent variable z with reparametrization
z = mu + torch.exp(log_sigma) * eps
reconstruction = self.decoder(z) # decode z
##### END TODO
#####

return {'q': q,
        'rec': reconstruction}

def loss(self, x, outputs):
    ##### TODO
    #####
    # Implement the loss computation of the VAE.
    #
    # HINT: Your code should implement the following steps:
    #
    #         1. compute the image reconstruction loss, similar to AE loss above
    #
    #         2. compute the KL divergence loss between the inferred posterior
    #
    #            distribution and a unit Gaussian prior; you can use the provided
    #
    #            function above for computing the KL divergence between two
Gaussians #
    #            parametrized by mean and log_sigma
    #
    # HINT: Make sure to compute the KL divergence in the correct order since it is
    #
    #         not symmetric!! ie.  $KL(p, q) \neq KL(q, p)$ 
    #
    #####
    loss_func = nn.BCELoss(reduction='mean')
    # loss_func = nn.MSELoss()
    # x: [64, 784], y: [64, 1, 784] -> [64, 784]
    rec_loss = loss_func(outputs['rec'].view(batch_size, -1), x)

    # print(outputs['q'].shape)
    q = outputs['q']
    # mu1 = q[:, :self.nz] # [batch_size, nz] = [64, 32]
    # log_sigma1 = q[:, self.nz:]
    mu1, log_sigma1 = torch.chunk(q, 2, dim=1)
    mu2 = torch.zeros_like(mu1).to(device)

```



```

        log_sigma2 = torch.zeros_like(log_sigma1).to(device)
        # make it as an scalar instead of a array
        # kl_loss = torch.mean(torch.sum(kl_divergence(mu1, log_sigma1, mu2,
log_sigma2), dim=1), dim=0)
        kl_loss = kl_divergence(mu1, log_sigma1, mu2,
log_sigma2).sum(dim=1).mean(dim=0)
        ##### END TODO
#####

        # return weighted objective
        return rec_loss + self.beta * kl_loss, \
            {'rec_loss': rec_loss, 'kl_loss': kl_loss}

def reconstruct(self, x):
    """Use mean of posterior estimate for visualization reconstruction."""
    ##### TODO
#####
    # This function is used for visualizing reconstructions of our VAE model. To
    #
    # obtain the maximum likelihood estimate we bypass the sampling procedure of
the #
    # inferred latent and instead directly use the mean of the inferred posterior.
    #
    # HINT: encode the input image and then decode the mean of the posterior to
obtain #
    #         the reconstruction.
    #

#####
    mu = self.encoder(x)[:self.nz]
    # encoder x, split mean, then decode mean
    image = self.decoder(mu).reshape(-1, 28, 28)
    ##### END TODO
#####
    return image

```

## Setting up the VAE Training Loop [4pt]

Let's start training the VAE model! We will first verify our implementation by setting  $\beta = 0$ .

```

# Prob1-5 VAE training loop
learning_rate = 1e-3
nz = 32
beta = 0

##### TODO #####
epochs = 5      # recommended 5-20 epochs
##### END TODO #####

```

```

# build VAE model
vae_model = VAE(nz, beta).to(device)    # transfer model to GPU if available
vae_model = vae_model.train()    # set model in train mode (eg batchnorm params get
updated)

# build optimizer and loss function
##### TODO #####
# Build the optimizer for the vae_model. We will again use the Adam optimizer with #
# the given learning rate and otherwise default parameters.                        #
#####
# same as AE
optimizer = torch.optim.Adam(vae_model.parameters(), lr=learning_rate)
##### END TODO #####

train_it = 0
rec_loss, kl_loss = [], []
print(f"Running {epochs} epochs with {beta=}")
for ep in range(epochs):
    print("Run Epoch {}".format(ep))
    ##### TODO
#####
    # Implement the main training loop for the VAE model.
#
# HINT: Your training loop should sample batches from the data loader, run the
#
#     forward pass of the VAE, compute the loss, perform the backward pass and
#
#     perform one gradient step with the optimizer.
#
# HINT: Don't forget to erase old gradients before performing the backward pass.
#
# HINT: This time we will use the loss() function of our model for computing the
#
#     training loss. It outputs the total training loss and a dict containing
#
#     the breakdown of reconstruction and KL loss.
#

#####
for X_train, y_train in mnist_data_loader:
    X_train, y_train = X_train.to(device), y_train.to(device)
    optimizer.zero_grad()
    X_train = X_train.reshape([batch_size, in_size])
    # falten X_train: [64, 28, 28] -> [64, 784]
    y_pred = vae_model(X_train)
    # print(y_pred['rec'].shape, X_train.shape)
    total_loss, losses = vae_model.loss(X_train, y_pred)
    # losses['rec_loss'] = losses['rec_loss'].detach().cpu()

```

```

# losses['kl_loss'] = losses['kl_loss'].detach().cpu()
# print(total_loss.shape)

total_loss.backward()
optimizer.step()

rec_loss.append(losses['rec_loss'])
kl_loss.append(losses['kl_loss'])
if train_it % 100 == 0:
    print("It {}: Total Loss: {}, \t Rec Loss: {},\t KL Loss: {}"
          .format(train_it, total_loss, losses['rec_loss'], losses['kl_loss']))
    train_it += 1
##### END TODO #####

print("Done!")

rec_loss_plotdata = [foo.detach().cpu() for foo in rec_loss]
kl_loss_plotdata = [foo.detach().cpu() for foo in kl_loss]

# log the loss training curves
fig = plt.figure(figsize = (10, 5))
ax1 = plt.subplot(121)
ax1.plot(rec_loss_plotdata)
ax1.title.set_text("Reconstruction Loss")
ax2 = plt.subplot(122)
ax2.plot(kl_loss_plotdata)
ax2.title.set_text("KL Loss")
plt.show()

```

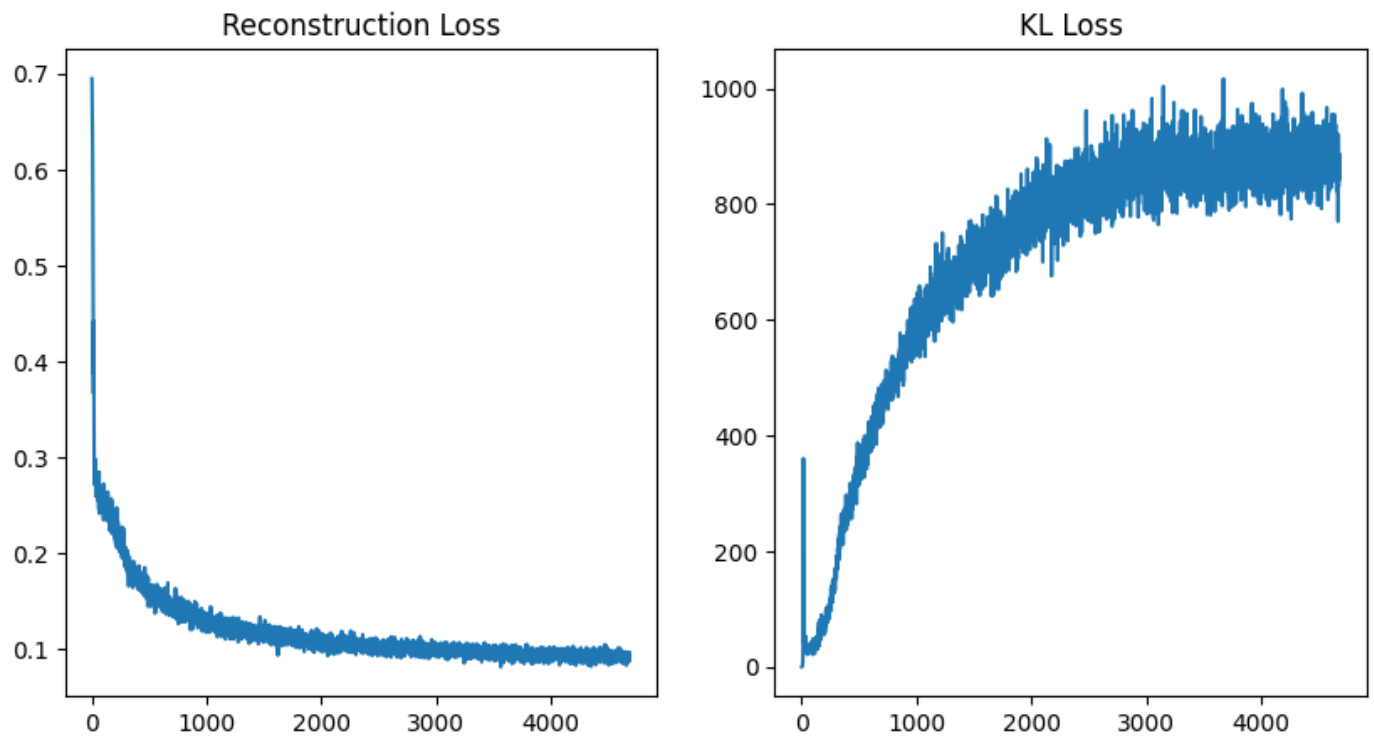
Running 5 epochs with beta=0

Run Epoch 0

It 0: Total Loss: 0.6949584484100342,	Rec Loss: 0.6949584484100342,	KL Loss: 0.30301231145858765
It 100: Total Loss: 0.25568607449531555,	Rec Loss: 0.25568607449531555,	KL Loss: 34.07611846923828
It 200: Total Loss: 0.236419677734375,	Rec Loss: 0.236419677734375,	KL Loss: 66.10553741455078
It 300: Total Loss: 0.20365391671657562,	Rec Loss: 0.20365391671657562,	KL Loss: 167.48959350585938
It 400: Total Loss: 0.16944284737110138,	Rec Loss: 0.16944284737110138,	KL Loss: 268.63836669921875
It 500: Total Loss: 0.16152262687683105,	Rec Loss: 0.16152262687683105,	KL Loss: 330.0137634277344
It 600: Total Loss: 0.15133586525917053,	Rec Loss: 0.15133586525917053,	KL Loss: 409.5831604003906
It 700: Total Loss: 0.147485613822937,	Rec Loss: 0.147485613822937,	KL Loss: 427.9168701171875
It 800: Total Loss: 0.14304688572883606,	Rec Loss: 0.14304688572883606,	KL Loss: 481.8742980957031

It 900: Total Loss: 0.13400740921497345, 537.3521728515625	Rec Loss: 0.13400740921497345,	KL Loss:
Run Epoch 1		
It 1000: Total Loss: 0.13504794239997864, 589.2233276367188	Rec Loss: 0.13504794239997864,	KL Loss:
It 1100: Total Loss: 0.11404438316822052, 637.029052734375	Rec Loss: 0.11404438316822052,	KL Loss:
It 1200: Total Loss: 0.11407783627510071, 616.8056030273438	Rec Loss: 0.11407783627510071,	KL Loss:
It 1300: Total Loss: 0.12522414326667786, 671.4837036132812	Rec Loss: 0.12522414326667786,	KL Loss:
It 1400: Total Loss: 0.11785966157913208, 676.5065307617188	Rec Loss: 0.11785966157913208,	KL Loss:
It 1500: Total Loss: 0.11737716943025589, 728.6448974609375	Rec Loss: 0.11737716943025589,	KL Loss:
It 1600: Total Loss: 0.11114905774593353, 718.3271484375	Rec Loss: 0.11114905774593353,	KL Loss:
It 1700: Total Loss: 0.11836127936840057, 699.9386596679688	Rec Loss: 0.11836127936840057,	KL Loss:
It 1800: Total Loss: 0.12070836871862411, 734.30224609375	Rec Loss: 0.12070836871862411,	KL Loss:
Run Epoch 2		
It 1900: Total Loss: 0.10240508615970612, 776.1109619140625	Rec Loss: 0.10240508615970612,	KL Loss:
It 2000: Total Loss: 0.10611988604068756, 818.9705810546875	Rec Loss: 0.10611988604068756,	KL Loss:
It 2100: Total Loss: 0.10807178914546967, 837.394775390625	Rec Loss: 0.10807178914546967,	KL Loss:
It 2200: Total Loss: 0.10127409547567368, 838.2537231445312	Rec Loss: 0.10127409547567368,	KL Loss:
It 2300: Total Loss: 0.1084408089518547, 813.2244873046875	Rec Loss: 0.1084408089518547,	KL Loss:
It 2400: Total Loss: 0.10154508799314499, 779.3074951171875	Rec Loss: 0.10154508799314499,	KL Loss:
It 2500: Total Loss: 0.10442092269659042, 850.5172119140625	Rec Loss: 0.10442092269659042,	KL Loss:
It 2600: Total Loss: 0.1011001393198967, 791.435302734375	Rec Loss: 0.1011001393198967,	KL Loss:
It 2700: Total Loss: 0.0972280353307724, 817.681396484375	Rec Loss: 0.0972280353307724,	KL Loss:
It 2800: Total Loss: 0.09908905625343323, 875.5213012695312	Rec Loss: 0.09908905625343323,	KL Loss:
Run Epoch 3		
It 2900: Total Loss: 0.10294036567211151, 882.3240356445312	Rec Loss: 0.10294036567211151,	KL Loss:
It 3000: Total Loss: 0.09609825909137726, 857.989013671875	Rec Loss: 0.09609825909137726,	KL Loss:
It 3100: Total Loss: 0.09500861167907715, 912.4680786132812	Rec Loss: 0.09500861167907715,	KL Loss:


It 3200: Total Loss: 0.09884800016880035, 897.725830078125	Rec Loss: 0.09884800016880035, KL Loss:
It 3300: Total Loss: 0.10174050182104111, 846.115478515625	Rec Loss: 0.10174050182104111, KL Loss:
It 3400: Total Loss: 0.09637000411748886, 898.1522216796875	Rec Loss: 0.09637000411748886, KL Loss:
It 3500: Total Loss: 0.09257102012634277, 826.9158325195312	Rec Loss: 0.09257102012634277, KL Loss:
It 3600: Total Loss: 0.08987121284008026, 824.9658203125	Rec Loss: 0.08987121284008026, KL Loss:
It 3700: Total Loss: 0.09394235163927078, 841.6760864257812	Rec Loss: 0.09394235163927078, KL Loss:
Run Epoch 4	
It 3800: Total Loss: 0.09188032150268555, 938.658203125	Rec Loss: 0.09188032150268555, KL Loss:
It 3900: Total Loss: 0.09344063699245453, 872.4083862304688	Rec Loss: 0.09344063699245453, KL Loss:
It 4000: Total Loss: 0.10003844648599625, 875.1181640625	Rec Loss: 0.10003844648599625, KL Loss:
It 4100: Total Loss: 0.09494234621524811, 863.547607421875	Rec Loss: 0.09494234621524811, KL Loss:
It 4200: Total Loss: 0.09472823143005371, 909.6565551757812	Rec Loss: 0.09472823143005371, KL Loss:
It 4300: Total Loss: 0.08830751478672028, 878.493408203125	Rec Loss: 0.08830751478672028, KL Loss:
It 4400: Total Loss: 0.09269534796476364, 855.6065673828125	Rec Loss: 0.09269534796476364, KL Loss:
It 4500: Total Loss: 0.08940350264310837, 824.0811767578125	Rec Loss: 0.08940350264310837, KL Loss:
It 4600: Total Loss: 0.10150223225355148, 834.7205200195312	Rec Loss: 0.10150223225355148, KL Loss:
Done!	




Let's look at some reconstructions and decoded embedding samples!

```
# visualize VAE reconstructions and samples from the generative model
print("beta = ", beta)
vis_reconstruction(vae_model, randomize=True)
vis_samples(vae_model)
```

```
beta = 0
```





## Tweaking the loss function $\beta$ [2pt]

Prob1-6: Let's repeat the same experiment for  $\beta = 10$ , a very high value for the coefficient.

```
# VAE training loop
learning_rate = 1e-3
nz = 32
beta = 10

##### TODO #####
epochs = 5      # recommended 5-20 epochs
##### END TODO #####
```

```

# build VAE model
vae_model = VAE(nz, beta).to(device)    # transfer model to GPU if available
vae_model = vae_model.train()    # set model in train mode (eg batchnorm params get
updated)

# build optimizer and loss function
##### TODO #####
# Build the optimizer for the vae_model. We will again use the Adam optimizer with #
# the given learning rate and otherwise default parameters.                        #
#####
# same as AE
optimizer = torch.optim.Adam(vae_model.parameters(), lr=learning_rate)
##### END TODO #####

train_it = 0
rec_loss, kl_loss = [], []
print(f"Running {epochs} epochs with {beta=}")
for ep in range(epochs):
    print("Run Epoch {}".format(ep))
    ##### TODO
    #####
    # Implement the main training loop for the VAE model.
    #
    # HINT: Your training loop should sample batches from the data loader, run the
    #
    #         forward pass of the VAE, compute the loss, perform the backward pass and
    #
    #         perform one gradient step with the optimizer.
    #
    # HINT: Don't forget to erase old gradients before performing the backward pass.
    #
    # HINT: This time we will use the loss() function of our model for computing the
    #
    #         training loss. It outputs the total training loss and a dict containing
    #
    #         the breakdown of reconstruction and KL loss.
    #

    #####
    for X_train, y_train in mnist_data_loader:
        X_train, y_train = X_train.to(device), y_train.to(device)
        optimizer.zero_grad()
        X_train = X_train.reshape([batch_size, in_size])
        # falten X_train: [64, 28, 28] -> [64, 784]
        y_pred = vae_model(X_train)
        # print(y_pred['rec'].shape, X_train.shape)
        total_loss, losses = vae_model.loss(X_train, y_pred)
        # losses['rec_loss'] = losses['rec_loss'].detach().cpu()
        # losses['kl_loss'] = losses['kl_loss'].detach().cpu()

```



```

# print(total_loss.shape)

total_loss.backward()
optimizer.step()

rec_loss.append(losses['rec_loss'])
kl_loss.append(losses['kl_loss'])
if train_it % 100 == 0:
    print("It {}: Total Loss: {}, \t Rec Loss: {},\t KL Loss: {}".format(
        train_it, total_loss, losses['rec_loss'], losses['kl_loss']))
    train_it += 1
##### END TODO #####

print("Done!")

rec_loss_plotdata = [foo.detach().cpu() for foo in rec_loss]
kl_loss_plotdata = [foo.detach().cpu() for foo in kl_loss]

# log the loss training curves
fig = plt.figure(figsize = (10, 5))
ax1 = plt.subplot(121)
ax1.plot(rec_loss_plotdata)
ax1.title.set_text("Reconstruction Loss")
ax2 = plt.subplot(122)
ax2.plot(kl_loss_plotdata)
ax2.title.set_text("KL Loss")
plt.show()

```

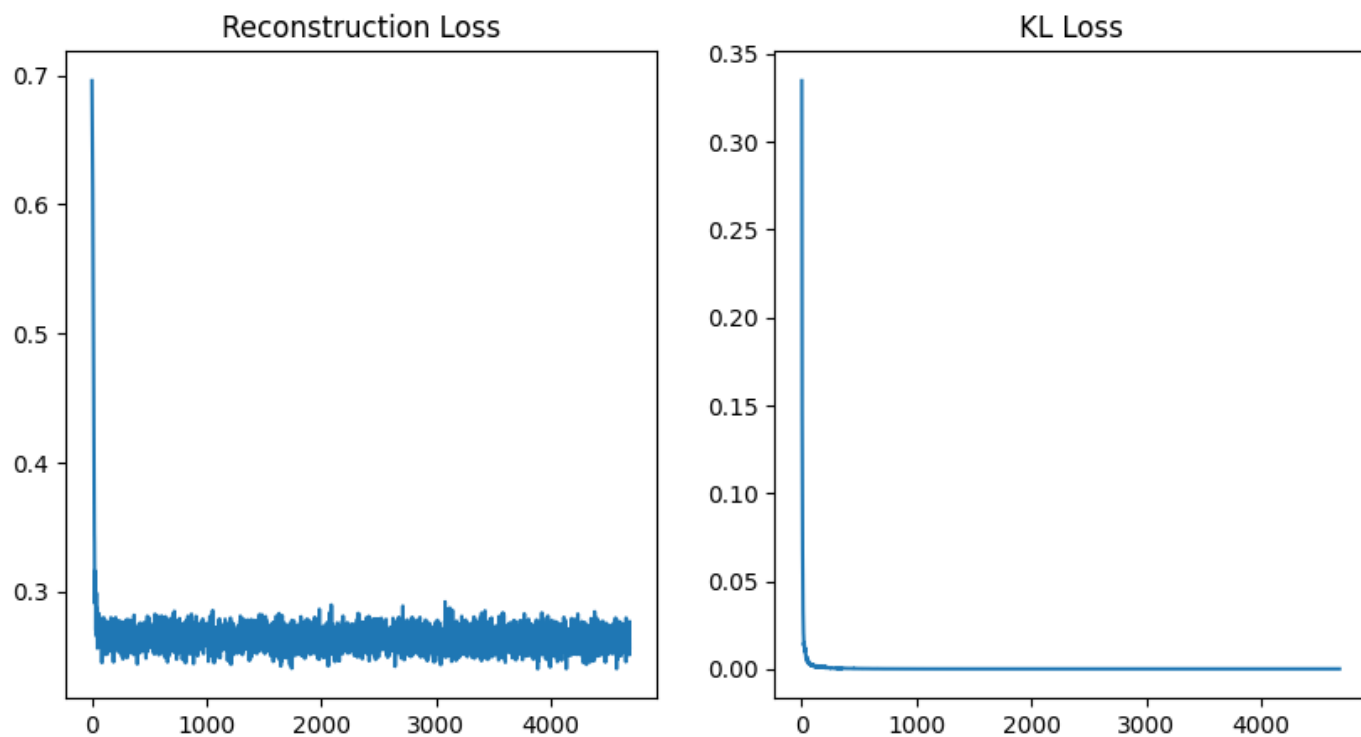
Running 5 epochs with beta=10

Run Epoch 0

It 0: Total Loss: 4.043502330780029, 0.3347723186016083	Rec Loss: 0.6957789659500122,	KL Loss:
It 100: Total Loss: 0.2881004810333252, 0.0017881318926811218	Rec Loss: 0.2702191472053528,	KL Loss:
It 200: Total Loss: 0.2783050537109375, 0.0011024139821529388	Rec Loss: 0.2672809064388275,	KL Loss:
It 300: Total Loss: 0.26480501890182495, 0.0003879130817949772	Rec Loss: 0.26092588901519775,	KL Loss:
It 400: Total Loss: 0.2728104293346405, 0.0003256821073591709	Rec Loss: 0.26955360174179077,	KL Loss:
It 500: Total Loss: 0.2586118280887604, 0.0001949113793671131	Rec Loss: 0.2566627264022827,	KL Loss:
It 600: Total Loss: 0.2576855421066284, 0.00012712227180600166	Rec Loss: 0.2564143240451813,	KL Loss:
It 700: Total Loss: 0.2710407078266144, 7.832655683159828e-05	Rec Loss: 0.27025744318962097,	KL Loss:
It 800: Total Loss: 0.2611941695213318, 6.949156522750854e-05	Rec Loss: 0.2604992389678955,	KL Loss:


It 900: Total Loss: 0.2668711841106415, 4.298286512494087e-05	Rec Loss: 0.26644134521484375,	KL Loss:
Run Epoch 1		
It 1000: Total Loss: 0.27138444781303406, 2.9797665774822235e-05	Rec Loss: 0.2710864841938019,	KL Loss:
It 1100: Total Loss: 0.2624884843826294, 2.599647268652916e-05	Rec Loss: 0.26222851872444153,	KL Loss:
It 1200: Total Loss: 0.26039522886276245, 1.6470439732074738e-05	Rec Loss: 0.26023051142692566,	KL Loss:
It 1300: Total Loss: 0.2614837884902954, 1.3169366866350174e-05	Rec Loss: 0.2613520920276642,	KL Loss:
It 1400: Total Loss: 0.2620086967945099, 1.0164454579353333e-05	Rec Loss: 0.26190704107284546,	KL Loss:
It 1500: Total Loss: 0.27102693915367126, 1.0045245289802551e-05	Rec Loss: 0.27092647552490234,	KL Loss:
It 1600: Total Loss: 0.2748018205165863, 7.689464837312698e-06	Rec Loss: 0.27472493052482605,	KL Loss:
It 1700: Total Loss: 0.26623067259788513, 8.478760719299316e-06	Rec Loss: 0.26614588499069214,	KL Loss:
It 1800: Total Loss: 0.26382920145988464, 4.9658119678497314e-06	Rec Loss: 0.26377955079078674,	KL Loss:
Run Epoch 2		
It 1900: Total Loss: 0.25132864713668823, 4.280358552932739e-06	Rec Loss: 0.2512858510017395,	KL Loss:
It 2000: Total Loss: 0.25856828689575195, 3.9138831198215485e-06	Rec Loss: 0.2585291564464569,	KL Loss:
It 2100: Total Loss: 0.25332406163215637, 3.180932253599167e-06	Rec Loss: 0.2532922625541687,	KL Loss:
It 2200: Total Loss: 0.27400606870651245, 3.1813979148864746e-06	Rec Loss: 0.2739742696285248,	KL Loss:
It 2300: Total Loss: 0.2571167051792145, 3.6847777664661407e-06	Rec Loss: 0.2570798695087433,	KL Loss:
It 2400: Total Loss: 0.26613470911979675, 3.259163349866867e-06	Rec Loss: 0.2661021053791046,	KL Loss:
It 2500: Total Loss: 0.27010685205459595, 2.671964466571808e-06	Rec Loss: 0.2700801193714142,	KL Loss:
It 2600: Total Loss: 0.2720736265182495, 2.6668421924114227e-06	Rec Loss: 0.2720469534397125,	KL Loss:
It 2700: Total Loss: 0.2632483243942261, 5.198176950216293e-06	Rec Loss: 0.26319634914398193,	KL Loss:
It 2800: Total Loss: 0.2606072425842285, 3.636348992586136e-06	Rec Loss: 0.2605708837509155,	KL Loss:
Run Epoch 3		
It 2900: Total Loss: 0.27058422565460205, 1.655425876379013e-06	Rec Loss: 0.2705676853656769,	KL Loss:
It 3000: Total Loss: 0.2606358528137207, 2.096407115459442e-06	Rec Loss: 0.26061490178108215,	KL Loss:
It 3100: Total Loss: 0.2866145968437195, 1.191161572933197e-06	Rec Loss: 0.2866026759147644,	KL Loss:


It 3200: Total Loss: 0.2566582262516022, 2.1960586309432983e-06	Rec Loss: 0.25663626194000244,	KL Loss:
It 3300: Total Loss: 0.275623083114624, 2.291519194841385e-06	Rec Loss: 0.2756001651287079,	KL Loss:
It 3400: Total Loss: 0.26475316286087036, 3.884546458721161e-06	Rec Loss: 0.2647143304347992,	KL Loss:
It 3500: Total Loss: 0.2548610270023346, 1.6363337635993958e-06	Rec Loss: 0.25484466552734375,	KL Loss:
It 3600: Total Loss: 0.26295459270477295, 3.661029040813446e-06	Rec Loss: 0.26291799545288086,	KL Loss:
It 3700: Total Loss: 0.2668231725692749, 1.2042000889778137e-06	Rec Loss: 0.2668111324310303,	KL Loss:
Run Epoch 4		
It 3800: Total Loss: 0.2698535621166229, 1.2186355888843536e-06	Rec Loss: 0.26984137296676636,	KL Loss:
It 3900: Total Loss: 0.2766888439655304, 1.4961697161197662e-06	Rec Loss: 0.2766738831996918,	KL Loss:
It 4000: Total Loss: 0.2591097354888916, 8.121132850646973e-07	Rec Loss: 0.25910162925720215,	KL Loss:
It 4100: Total Loss: 0.2572685778141022, 8.246861398220062e-07	Rec Loss: 0.2572603225708008,	KL Loss:
It 4200: Total Loss: 0.2516791820526123, 9.583309292793274e-07	Rec Loss: 0.25166958570480347,	KL Loss:
It 4300: Total Loss: 0.2628110349178314, 1.1031515896320343e-06	Rec Loss: 0.262800008058548,	KL Loss:
It 4400: Total Loss: 0.2500588595867157, 9.918585419654846e-07	Rec Loss: 0.2500489354133606,	KL Loss:
It 4500: Total Loss: 0.27144747972488403, 1.991167664527893e-06	Rec Loss: 0.27142757177352905,	KL Loss:
It 4600: Total Loss: 0.26362940669059753, 1.2149102985858917e-06	Rec Loss: 0.26361724734306335,	KL Loss:
Done!		



```
# visualize VAE reconstructions and samples from the generative model
print("beta = ", beta)
vis_reconstruction(vae_model, randomize=True)
vis_samples(vae_model)
```

```
beta = 10
```





**Inline Question: What can you observe when setting  $\beta = 0$  and  $\beta = 10$ ? Explain your observations! [2pt]**

(max 200 words)

**Answer:**

When setting  $\beta = 0$  in VAE, the KL divergence term in the loss function is completely removed, and the model only considers the reconstruction loss. In this case, the VAE model behaves like an AE model, and is without the regularization term  $\beta$  on the latent space distribution. From the figure, the reconstruction image is the same as the AE model before. Therefore, the VAE with  $\beta = 0$  doesn't have a good generative process. And after decoding, sample distribution cannot produce clear results.

When setting  $\beta = 10$  in VAE, the weight of the KL divergence term is heavily increased. This encourages the VAE model to make the latent space more continuous and smooth, reducing the difference between the reconstructed images and the latent variable. If  $\beta > 0$ , a stronger constraint will be applied to latent space. However, if  $\beta$  is set too large (like 10), the VAE model may become over-constrained, it make too strong constraint at the cost of reconstruction quality. As we can see in the figure, the reconstructions and the samples are blurred. If we constrain a very low KL loss, the VAE's reconstructed quality will also decreases

Thus, it's essential for us to make trade-off between latent space amd reconstructions. Choosing an appropriate value of  $\beta$  can help the model better balance the weight of reconstruction loss and KL divergence term, leading to better latent variable representation and generative ability.

## Obtaining the best $\beta$ -factor [5pt]

Prob 1-6 continued: Now we can start tuning the beta value to achieve a good result. First describe what a "good result" would look like (focus what you would expect for reconstructions and sample quality).

**Inline Question: Characterize what properties you would expect for reconstructions and samples of a well-tuned VAE! [3pt]**

(max 200 words)

**Answer:**

A well-tuned VAE is expected to have:

1. High reconstructions quality. The reconstruced images shoud look like the original images with low blurred level, i.e., The VAE should be able to accurately reconstruct the input data from the latent space representation.
2. Low reconstruction error. The reconstruction error, i.e., the difference between the input data and its reconstruction, should be low.
3. Consistency between reconstructions & samples. The reconstructed images and the generated samples should be consistent with the learned data distribution. This indicates that the VAE is able to capture the underlying data distribution accurately.
4. Smooth and continuous latent space representation. Neighboring points in the latent space are corresponding to similar input data points. Compared to AE, VAE can generate high-quality samples from the learned latent space distribution.

Now that you know what outcome we would like to obtain, try to tune  $\beta$  to achieve this result. Logarithmic search in steps of 10x will be helpful, good results can be achieved after ~20 epochs of training. Training reconstructions should be high quality, test samples should be diverse, distinguishable numbers, most samples recognizable as numbers.

**Answer: Tuned beta value 0.005 [2pt]**

```
# Tuning for best beta
learning_rate = 1e-3
nz = 32

##### TODO #####
epochs = 20      # recommended 5-20 epochs
beta = 0.005 # Tune this for best results
```

```
##### END TODO #####

# build VAE model
vae_model = VAE(nz, beta).to(device)    # transfer model to GPU if available
vae_model = vae_model.train()    # set model in train mode (eg batchnorm params get
updated)

# build optimizer and loss function
##### TODO #####
# Build the optimizer for the vae_model. We will again use the Adam optimizer with #
# the given learning rate and otherwise default parameters.                        #
#####
# same as AE
optimizer = torch.optim.Adam(vae_model.parameters(), lr=learning_rate)
##### END TODO #####

train_it = 0
rec_loss, kl_loss = [], []
print(f"Running {epochs} epochs with {beta=}")
for ep in range(epochs):
    print("Run Epoch {}".format(ep))
    ##### TODO #####
#####
    # Implement the main training loop for the VAE model.
#
# HINT: Your training loop should sample batches from the data loader, run the
#
#     forward pass of the VAE, compute the loss, perform the backward pass and
#
#     perform one gradient step with the optimizer.
#
# HINT: Don't forget to erase old gradients before performing the backward pass.
#
# HINT: This time we will use the loss() function of our model for computing the
#
#     training loss. It outputs the total training loss and a dict containing
#
#     the breakdown of reconstruction and KL loss.
#
#####
    for X_train, y_train in mnist_data_loader:
        X_train, y_train = X_train.to(device), y_train.to(device)
        optimizer.zero_grad()
        X_train = X_train.reshape([batch_size, in_size])
        # falten X_train: [64, 28, 28] -> [64, 784]
        y_pred = vae_model(X_train)
        # print(y_pred['rec'].shape, X_train.shape)
        total_loss, losses = vae_model.loss(X_train, y_pred)
```

```

# losses['rec_loss'] = losses['rec_loss'].detach().cpu()
# losses['kl_loss'] = losses['kl_loss'].detach().cpu()
# print(total_loss.shape)

total_loss.backward()
optimizer.step()

rec_loss.append(losses['rec_loss'])
kl_loss.append(losses['kl_loss'])
if train_it % 100 == 0:
    print("It {}: Total Loss: {}, \t Rec Loss: {},\t KL Loss: {}"
          .format(train_it, total_loss, losses['rec_loss'], losses['kl_loss']))
    train_it += 1
##### END TODO #####

print("Done!")

rec_loss_plotdata = [foo.detach().cpu() for foo in rec_loss]
kl_loss_plotdata = [foo.detach().cpu() for foo in kl_loss]

# log the loss training curves
fig = plt.figure(figsize = (10, 5))
ax1 = plt.subplot(121)
ax1.plot(rec_loss_plotdata)
ax1.title.set_text("Reconstruction Loss")
ax2 = plt.subplot(122)
ax2.plot(kl_loss_plotdata)
ax2.title.set_text("KL Loss")
plt.show()

```

Running 20 epochs with beta=0.005

Run Epoch 0

It 0: Total Loss: 0.6951301097869873,	Rec Loss: 0.6934783458709717,	KL Loss: 0.33035051822662354
It 100: Total Loss: 0.2443913072347641,	Rec Loss: 0.24127808213233948,	KL Loss: 0.6226447224617004
It 200: Total Loss: 0.24876989424228668,	Rec Loss: 0.24097298085689545,	KL Loss: 1.559383749961853
It 300: Total Loss: 0.24794985353946686,	Rec Loss: 0.24170154333114624,	KL Loss: 1.2496623992919922
It 400: Total Loss: 0.2551601231098175,	Rec Loss: 0.24916376173496246,	KL Loss: 1.1992732286453247
It 500: Total Loss: 0.2513655722141266,	Rec Loss: 0.24168802797794342,	KL Loss: 1.9355098009109497
It 600: Total Loss: 0.24870824813842773,	Rec Loss: 0.23745959997177124,	KL Loss: 2.249729633331299
It 700: Total Loss: 0.2450023591518402,	Rec Loss: 0.23487970232963562,	KL Loss: 2.0245306491851807



It 800: Total Loss: 0.24700921773910522, 2.895991325378418	Rec Loss: 0.23252926766872406,	KL Loss:
It 900: Total Loss: 0.2514171600341797, 3.3882479667663574	Rec Loss: 0.23447592556476593,	KL Loss:
Run Epoch 1		
It 1000: Total Loss: 0.22184057533740997, 3.4779553413391113	Rec Loss: 0.2044508010149002,	KL Loss:
It 1100: Total Loss: 0.24142208695411682, 4.164727687835693	Rec Loss: 0.22059844434261322,	KL Loss:
It 1200: Total Loss: 0.23203182220458984, 5.009239196777344	Rec Loss: 0.20698562264442444,	KL Loss:
It 1300: Total Loss: 0.22673968970775604, 5.0605902671813965	Rec Loss: 0.20143674314022064,	KL Loss:
It 1400: Total Loss: 0.2066623717546463, 5.208342552185059	Rec Loss: 0.18062065541744232,	KL Loss:
It 1500: Total Loss: 0.22205908596515656, 4.8580474853515625	Rec Loss: 0.19776885211467743,	KL Loss:
It 1600: Total Loss: 0.2141190469264984, 5.350385665893555	Rec Loss: 0.18736711144447327,	KL Loss:
It 1700: Total Loss: 0.21713028848171234, 5.240086078643799	Rec Loss: 0.1909298598766327,	KL Loss:
It 1800: Total Loss: 0.194880872964859, 5.496772766113281	Rec Loss: 0.16739700734615326,	KL Loss:
Run Epoch 2		
It 1900: Total Loss: 0.19650927186012268, 5.7187395095825195	Rec Loss: 0.16791556775569916,	KL Loss:
It 2000: Total Loss: 0.20267681777477264, 6.078100204467773	Rec Loss: 0.17228631675243378,	KL Loss:
It 2100: Total Loss: 0.19628649950027466, 5.68217658996582	Rec Loss: 0.16787561774253845,	KL Loss:
It 2200: Total Loss: 0.196811243891716, 5.796411514282227	Rec Loss: 0.16782918572425842,	KL Loss:
It 2300: Total Loss: 0.1915660798549652, 6.40994930267334	Rec Loss: 0.1595163345336914,	KL Loss:
It 2400: Total Loss: 0.1947268694639206, 5.786983489990234	Rec Loss: 0.16579195857048035,	KL Loss:
It 2500: Total Loss: 0.19760264456272125, 6.155581474304199	Rec Loss: 0.16682474315166473,	KL Loss:
It 2600: Total Loss: 0.203748419880867, 6.394340991973877	Rec Loss: 0.17177671194076538,	KL Loss:
It 2700: Total Loss: 0.2055603563785553, 6.659918308258057	Rec Loss: 0.17226076126098633,	KL Loss:
It 2800: Total Loss: 0.20353147387504578, 6.534791946411133	Rec Loss: 0.1708575189113617,	KL Loss:
Run Epoch 3		
It 2900: Total Loss: 0.1967518925666809, 6.265541076660156	Rec Loss: 0.165424183011055,	KL Loss:
It 3000: Total Loss: 0.2000911384820938, 6.830418586730957	Rec Loss: 0.16593904793262482,	KL Loss:

It 3100: Total Loss: 0.20262224972248077, 6.3062896728515625	Rec Loss: 0.17109079658985138, KL Loss:
It 3200: Total Loss: 0.19437962770462036, 6.413968086242676	Rec Loss: 0.16230979561805725, KL Loss:
It 3300: Total Loss: 0.21200965344905853, 6.8367180824279785	Rec Loss: 0.17782606184482574, KL Loss:
It 3400: Total Loss: 0.20770622789859772, 6.651199817657471	Rec Loss: 0.17445023357868195, KL Loss:
It 3500: Total Loss: 0.19349583983421326, 6.544774055480957	Rec Loss: 0.16077196598052979, KL Loss:
It 3600: Total Loss: 0.20763127505779266, 7.123852729797363	Rec Loss: 0.17201201617717743, KL Loss:
It 3700: Total Loss: 0.20930856466293335, 6.920116901397705	Rec Loss: 0.17470797896385193, KL Loss:
Run Epoch 4	
It 3800: Total Loss: 0.19993086159229279, 7.087825775146484	Rec Loss: 0.16449172794818878, KL Loss:
It 3900: Total Loss: 0.20996065437793732, 7.01833438873291	Rec Loss: 0.17486898601055145, KL Loss:
It 4000: Total Loss: 0.20386317372322083, 6.868167400360107	Rec Loss: 0.16952234506607056, KL Loss:
It 4100: Total Loss: 0.1966114193201065, 6.865027904510498	Rec Loss: 0.16228628158569336, KL Loss:
It 4200: Total Loss: 0.1865125596523285, 6.664539337158203	Rec Loss: 0.15318986773490906, KL Loss:
It 4300: Total Loss: 0.19731616973876953, 6.676242828369141	Rec Loss: 0.16393496096134186, KL Loss:
It 4400: Total Loss: 0.19014900922775269, 6.826122283935547	Rec Loss: 0.15601840615272522, KL Loss:
It 4500: Total Loss: 0.19030043482780457, 6.663074493408203	Rec Loss: 0.1569850593805313, KL Loss:
It 4600: Total Loss: 0.19880101084709167, 7.162077903747559	Rec Loss: 0.16299062967300415, KL Loss:
Run Epoch 5	
It 4700: Total Loss: 0.19488291442394257, 6.972219944000244	Rec Loss: 0.1600218117237091, KL Loss:
It 4800: Total Loss: 0.1990889012813568, 6.852155685424805	Rec Loss: 0.1648281216621399, KL Loss:
It 4900: Total Loss: 0.19571277499198914, 6.866893768310547	Rec Loss: 0.16137830913066864, KL Loss:
It 5000: Total Loss: 0.1889776885509491, 7.04770565032959	Rec Loss: 0.15373916923999786, KL Loss:
It 5100: Total Loss: 0.20280897617340088, 7.047110557556152	Rec Loss: 0.16757342219352722, KL Loss:
It 5200: Total Loss: 0.18687373399734497, 6.8358564376831055	Rec Loss: 0.1526944488286972, KL Loss:
It 5300: Total Loss: 0.20477578043937683, 7.093958854675293	Rec Loss: 0.16930599510669708, KL Loss:

It 5400: Total Loss: 0.19899554550647736, 7.359989643096924	Rec Loss: 0.1621955931186676,	KL Loss:
It 5500: Total Loss: 0.18705078959465027, 7.096412658691406	Rec Loss: 0.1515687257051468,	KL Loss:
It 5600: Total Loss: 0.18706177175045013, 6.844776630401611	Rec Loss: 0.15283788740634918,	KL Loss:
Run Epoch 6		
It 5700: Total Loss: 0.1935669332742691, 7.255548477172852	Rec Loss: 0.15728919208049774,	KL Loss:
It 5800: Total Loss: 0.1895630657672882, 7.326891899108887	Rec Loss: 0.15292860567569733,	KL Loss:
It 5900: Total Loss: 0.18698394298553467, 7.055345058441162	Rec Loss: 0.1517072170972824,	KL Loss:
It 6000: Total Loss: 0.18950411677360535, 7.030004978179932	Rec Loss: 0.15435409545898438,	KL Loss:
It 6100: Total Loss: 0.18968701362609863, 7.001849174499512	Rec Loss: 0.15467776358127594,	KL Loss:
It 6200: Total Loss: 0.19093656539916992, 7.540716171264648	Rec Loss: 0.15323297679424286,	KL Loss:
It 6300: Total Loss: 0.1920289546251297, 7.319461822509766	Rec Loss: 0.15543164312839508,	KL Loss:
It 6400: Total Loss: 0.19125127792358398, 7.199246406555176	Rec Loss: 0.1552550494670868,	KL Loss:
It 6500: Total Loss: 0.19423648715019226, 7.190062046051025	Rec Loss: 0.1582861691713333,	KL Loss:
Run Epoch 7		
It 6600: Total Loss: 0.18539518117904663, 7.252843856811523	Rec Loss: 0.1491309553384781,	KL Loss:
It 6700: Total Loss: 0.18688955903053284, 7.337984085083008	Rec Loss: 0.15019963681697845,	KL Loss:
It 6800: Total Loss: 0.18603327870368958, 7.022171497344971	Rec Loss: 0.15092241764068604,	KL Loss:
It 6900: Total Loss: 0.18600872159004211, 7.1846489906311035	Rec Loss: 0.1500854790210724,	KL Loss:
It 7000: Total Loss: 0.1882767677307129, 7.567228317260742	Rec Loss: 0.15044063329696655,	KL Loss:
It 7100: Total Loss: 0.1843862235546112, 7.626364707946777	Rec Loss: 0.14625440537929535,	KL Loss:
It 7200: Total Loss: 0.1825878918170929, 7.363788604736328	Rec Loss: 0.145768940448761,	KL Loss:
It 7300: Total Loss: 0.1900421679019928, 7.420465469360352	Rec Loss: 0.1529398411512375,	KL Loss:
It 7400: Total Loss: 0.19006752967834473, 7.433444499969482	Rec Loss: 0.1529003083705902,	KL Loss:
Run Epoch 8		
It 7500: Total Loss: 0.19654998183250427, 7.618408203125	Rec Loss: 0.1584579348564148,	KL Loss:
It 7600: Total Loss: 0.1849372833967209, 7.479524612426758	Rec Loss: 0.1475396603345871,	KL Loss:

It 7700: Total Loss: 0.18161731958389282, 7.419319152832031	Rec Loss: 0.14452072978019714,	KL Loss:
It 7800: Total Loss: 0.19282472133636475, 7.552432537078857	Rec Loss: 0.15506255626678467,	KL Loss:
It 7900: Total Loss: 0.19747202098369598, 7.375269412994385	Rec Loss: 0.16059567034244537,	KL Loss:
It 8000: Total Loss: 0.19068089127540588, 7.841041564941406	Rec Loss: 0.1514756828546524,	KL Loss:
It 8100: Total Loss: 0.19213226437568665, 7.718325614929199	Rec Loss: 0.15354064106941223,	KL Loss:
It 8200: Total Loss: 0.20216801762580872, 7.792954444885254	Rec Loss: 0.16320323944091797,	KL Loss:
It 8300: Total Loss: 0.17858733236789703, 7.3956427574157715	Rec Loss: 0.14160911738872528,	KL Loss:
It 8400: Total Loss: 0.18532830476760864, 7.900510787963867	Rec Loss: 0.14582575857639313,	KL Loss:
Run Epoch 9		
It 8500: Total Loss: 0.18583723902702332, 7.770216464996338	Rec Loss: 0.14698615670204163,	KL Loss:
It 8600: Total Loss: 0.1969892382621765, 7.7090559005737305	Rec Loss: 0.15844395756721497,	KL Loss:
It 8700: Total Loss: 0.19546076655387878, 7.486700057983398	Rec Loss: 0.1580272763967514,	KL Loss:
It 8800: Total Loss: 0.18930882215499878, 7.939485549926758	Rec Loss: 0.14961139857769012,	KL Loss:
It 8900: Total Loss: 0.18677033483982086, 8.105522155761719	Rec Loss: 0.14624272286891937,	KL Loss:
It 9000: Total Loss: 0.18644224107265472, 7.2733306884765625	Rec Loss: 0.15007558465003967,	KL Loss:
It 9100: Total Loss: 0.18947285413742065, 7.740213394165039	Rec Loss: 0.15077179670333862,	KL Loss:
It 9200: Total Loss: 0.1850229799747467, 7.7744221687316895	Rec Loss: 0.1461508721113205,	KL Loss:
It 9300: Total Loss: 0.19828584790229797, 7.572961807250977	Rec Loss: 0.16042104363441467,	KL Loss:
Run Epoch 10		
It 9400: Total Loss: 0.18824121356010437, 7.855619430541992	Rec Loss: 0.14896312355995178,	KL Loss:
It 9500: Total Loss: 0.18680152297019958, 7.143072128295898	Rec Loss: 0.15108616650104523,	KL Loss:
It 9600: Total Loss: 0.18067589402198792, 7.3112592697143555	Rec Loss: 0.14411959052085876,	KL Loss:
It 9700: Total Loss: 0.18653658032417297, 7.819928169250488	Rec Loss: 0.1474369317293167,	KL Loss:
It 9800: Total Loss: 0.18674400448799133, 7.591657638549805	Rec Loss: 0.14878571033477783,	KL Loss:
It 9900: Total Loss: 0.1909576803445816, 7.757437705993652	Rec Loss: 0.15217049419879913,	KL Loss:


It 10000: Total Loss: 0.18401560187339783, 7.9246344566345215	Rec Loss: 0.14439243078231812,	KL Loss:
It 10100: Total Loss: 0.18373335897922516, 7.456473350524902	Rec Loss: 0.14645099639892578,	KL Loss:
It 10200: Total Loss: 0.18784180283546448, 7.8855485916137695	Rec Loss: 0.14841406047344208,	KL Loss:
It 10300: Total Loss: 0.18322569131851196, 7.239912033081055	Rec Loss: 0.14702613651752472,	KL Loss:
Run Epoch 11		
It 10400: Total Loss: 0.18788081407546997, 7.816140174865723	Rec Loss: 0.1488001048564911,	KL Loss:
It 10500: Total Loss: 0.18667256832122803, 7.498013496398926	Rec Loss: 0.1491824984550476,	KL Loss:
It 10600: Total Loss: 0.19172543287277222, 7.344147205352783	Rec Loss: 0.15500469505786896,	KL Loss:
It 10700: Total Loss: 0.1872260868549347, 7.66586971282959	Rec Loss: 0.1488967388868332,	KL Loss:
It 10800: Total Loss: 0.17188550531864166, 7.4849419593811035	Rec Loss: 0.13446079194545746,	KL Loss:
It 10900: Total Loss: 0.19570723176002502, 7.86707067489624	Rec Loss: 0.15637187659740448,	KL Loss:
It 11000: Total Loss: 0.1935248076915741, 7.7421345710754395	Rec Loss: 0.15481413900852203,	KL Loss:
It 11100: Total Loss: 0.18665218353271484, 7.8788347244262695	Rec Loss: 0.14725801348686218,	KL Loss:
It 11200: Total Loss: 0.189425989985466, 7.615576267242432	Rec Loss: 0.15134811401367188,	KL Loss:
Run Epoch 12		
It 11300: Total Loss: 0.1935732662677765, 7.584600448608398	Rec Loss: 0.1556502729654312,	KL Loss:
It 11400: Total Loss: 0.19450336694717407, 7.970778942108154	Rec Loss: 0.15464948117733002,	KL Loss:
It 11500: Total Loss: 0.18322308361530304, 8.002958297729492	Rec Loss: 0.14320829510688782,	KL Loss:
It 11600: Total Loss: 0.1875043362379074, 7.666605472564697	Rec Loss: 0.14917130768299103,	KL Loss:
It 11700: Total Loss: 0.19429203867912292, 8.237565994262695	Rec Loss: 0.15310420095920563,	KL Loss:
It 11800: Total Loss: 0.1964062750339508, 7.7469305992126465	Rec Loss: 0.15767161548137665,	KL Loss:
It 11900: Total Loss: 0.18954935669898987, 7.6480913162231445	Rec Loss: 0.15130889415740967,	KL Loss:
It 12000: Total Loss: 0.19058331847190857, 8.070592880249023	Rec Loss: 0.15023036301136017,	KL Loss:
It 12100: Total Loss: 0.1869863122701645, 7.656482696533203	Rec Loss: 0.1487039029598236,	KL Loss:
Run Epoch 13		
It 12200: Total Loss: 0.17617058753967285, 7.791522979736328	Rec Loss: 0.13721297681331635,	KL Loss:

It 12300: Total Loss: 0.18486186861991882, 7.753015518188477	Rec Loss: 0.14609679579734802,	KL Loss:
It 12400: Total Loss: 0.2003490924835205, 7.71613883972168	Rec Loss: 0.16176839172840118,	KL Loss:
It 12500: Total Loss: 0.19093936681747437, 7.313527584075928	Rec Loss: 0.1543717384338379,	KL Loss:
It 12600: Total Loss: 0.19128888845443726, 7.824732780456543	Rec Loss: 0.1521652191877365,	KL Loss:
It 12700: Total Loss: 0.190524622797966, 8.008378982543945	Rec Loss: 0.15048272907733917,	KL Loss:
It 12800: Total Loss: 0.18168240785598755, 7.7829718589782715	Rec Loss: 0.1427675485610962,	KL Loss:
It 12900: Total Loss: 0.18937534093856812, 7.649273872375488	Rec Loss: 0.15112897753715515,	KL Loss:
It 13000: Total Loss: 0.18815547227859497, 7.736161231994629	Rec Loss: 0.14947466552257538,	KL Loss:
It 13100: Total Loss: 0.19273564219474792, 7.994407653808594	Rec Loss: 0.15276360511779785,	KL Loss:
Run Epoch 14		
It 13200: Total Loss: 0.18535305559635162, 7.324501991271973	Rec Loss: 0.1487305462360382,	KL Loss:
It 13300: Total Loss: 0.183826744556427, 7.766269207000732	Rec Loss: 0.14499539136886597,	KL Loss:
It 13400: Total Loss: 0.18496474623680115, 7.917668342590332	Rec Loss: 0.14537641406059265,	KL Loss:
It 13500: Total Loss: 0.19278588891029358, 7.848118782043457	Rec Loss: 0.1535452902317047,	KL Loss:
It 13600: Total Loss: 0.18447203934192657, 7.440271377563477	Rec Loss: 0.14727067947387695,	KL Loss:
It 13700: Total Loss: 0.18202924728393555, 7.881281852722168	Rec Loss: 0.14262284338474274,	KL Loss:
It 13800: Total Loss: 0.18698722124099731, 8.18058967590332	Rec Loss: 0.1460842788219452,	KL Loss:
It 13900: Total Loss: 0.18116974830627441, 7.393972396850586	Rec Loss: 0.1441998928785324,	KL Loss:
It 14000: Total Loss: 0.18770824372768402, 7.937355995178223	Rec Loss: 0.14802145957946777,	KL Loss:
Run Epoch 15		
It 14100: Total Loss: 0.1852949857711792, 7.792640209197998	Rec Loss: 0.146331787109375,	KL Loss:
It 14200: Total Loss: 0.1930757611989975, 8.088260650634766	Rec Loss: 0.15263445675373077,	KL Loss:
It 14300: Total Loss: 0.18235264718532562, 8.222996711730957	Rec Loss: 0.14123766124248505,	KL Loss:
It 14400: Total Loss: 0.19115066528320312, 7.936586380004883	Rec Loss: 0.1514677256345749,	KL Loss:
It 14500: Total Loss: 0.1819249540567398, 7.433035373687744	Rec Loss: 0.14475977420806885,	KL Loss:

It 14600: Total Loss: 0.18414416909217834, 7.9186787605285645	Rec Loss: 0.14455077052116394,	KL Loss:
It 14700: Total Loss: 0.18654689192771912, 7.678439140319824	Rec Loss: 0.14815469086170197,	KL Loss:
It 14800: Total Loss: 0.1933647245168686, 8.179423332214355	Rec Loss: 0.15246760845184326,	KL Loss:
It 14900: Total Loss: 0.19114679098129272, 8.156071662902832	Rec Loss: 0.1503664255142212,	KL Loss:
Run Epoch 16		
It 15000: Total Loss: 0.18181112408638, 8.129806518554688	Rec Loss: 0.14116209745407104,	KL Loss:
It 15100: Total Loss: 0.1729908287525177, 7.318188667297363	Rec Loss: 0.13639989495277405,	KL Loss:
It 15200: Total Loss: 0.1822843998670578, 8.052167892456055	Rec Loss: 0.14202356338500977,	KL Loss:
It 15300: Total Loss: 0.1919952630996704, 8.32178020477295	Rec Loss: 0.15038636326789856,	KL Loss:
It 15400: Total Loss: 0.17490413784980774, 8.243961334228516	Rec Loss: 0.13368433713912964,	KL Loss:
It 15500: Total Loss: 0.18474510312080383, 7.948615074157715	Rec Loss: 0.14500203728675842,	KL Loss:
It 15600: Total Loss: 0.18608561158180237, 7.860692977905273	Rec Loss: 0.14678214490413666,	KL Loss:
It 15700: Total Loss: 0.1844157725572586, 7.661670684814453	Rec Loss: 0.14610742032527924,	KL Loss:
It 15800: Total Loss: 0.18887898325920105, 7.866703987121582	Rec Loss: 0.14954546093940735,	KL Loss:
It 15900: Total Loss: 0.18216419219970703, 8.057074546813965	Rec Loss: 0.14187881350517273,	KL Loss:
Run Epoch 17		
It 16000: Total Loss: 0.18552523851394653, 8.149752616882324	Rec Loss: 0.14477647840976715,	KL Loss:
It 16100: Total Loss: 0.1823066622018814, 7.667329788208008	Rec Loss: 0.14397001266479492,	KL Loss:
It 16200: Total Loss: 0.19213439524173737, 8.385414123535156	Rec Loss: 0.15020732581615448,	KL Loss:
It 16300: Total Loss: 0.17049941420555115, 8.069877624511719	Rec Loss: 0.13015002012252808,	KL Loss:
It 16400: Total Loss: 0.18184056878089905, 7.7235894203186035	Rec Loss: 0.1432226300239563,	KL Loss:
It 16500: Total Loss: 0.1807950735092163, 7.7659831047058105	Rec Loss: 0.14196515083312988,	KL Loss:
It 16600: Total Loss: 0.191180020570755, 7.963287353515625	Rec Loss: 0.1513635814189911,	KL Loss:
It 16700: Total Loss: 0.1910543441772461, 8.186273574829102	Rec Loss: 0.1501229852437973,	KL Loss:
It 16800: Total Loss: 0.1856914460659027, 7.953529357910156	Rec Loss: 0.14592380821704865,	KL Loss:
Run Epoch 18		

It 16900: Total Loss: 0.1933879554271698, 7.905253887176514	Rec Loss: 0.15386168658733368,	KL Loss:
It 17000: Total Loss: 0.1883453130722046, 8.251703262329102	Rec Loss: 0.14708679914474487,	KL Loss:
It 17100: Total Loss: 0.19285333156585693, 8.434236526489258	Rec Loss: 0.15068215131759644,	KL Loss:
It 17200: Total Loss: 0.18289020657539368, 7.8108625411987305	Rec Loss: 0.1438358873128891,	KL Loss:
It 17300: Total Loss: 0.18287375569343567, 7.984746932983398	Rec Loss: 0.14295001327991486,	KL Loss:
It 17400: Total Loss: 0.1826770305633545, 7.996283531188965	Rec Loss: 0.1426956057548523,	KL Loss:
It 17500: Total Loss: 0.20027422904968262, 8.428141593933105	Rec Loss: 0.15813352167606354,	KL Loss:
It 17600: Total Loss: 0.18573778867721558, 7.480560302734375	Rec Loss: 0.14833498001098633,	KL Loss:
It 17700: Total Loss: 0.18565452098846436, 8.254549026489258	Rec Loss: 0.14438177645206451,	KL Loss:
It 17800: Total Loss: 0.1775231510400772, 7.802834510803223	Rec Loss: 0.13850897550582886,	KL Loss:
Run Epoch 19		
It 17900: Total Loss: 0.18767033517360687, 8.137656211853027	Rec Loss: 0.14698205888271332,	KL Loss:
It 18000: Total Loss: 0.1843641847372055, 7.604427337646484	Rec Loss: 0.14634205400943756,	KL Loss:
It 18100: Total Loss: 0.17245778441429138, 7.9465861320495605	Rec Loss: 0.1327248513698578,	KL Loss:
It 18200: Total Loss: 0.18513017892837524, 7.8517165184021	Rec Loss: 0.1458715945482254,	KL Loss:
It 18300: Total Loss: 0.1836102306842804, 7.983452320098877	Rec Loss: 0.1436929702758789,	KL Loss:
It 18400: Total Loss: 0.17901290953159332, 7.847348213195801	Rec Loss: 0.13977617025375366,	KL Loss:
It 18500: Total Loss: 0.18653157353401184, 7.982309341430664	Rec Loss: 0.1466200351715088,	KL Loss:
It 18600: Total Loss: 0.18460577726364136, 8.00623607635498	Rec Loss: 0.1445745974779129,	KL Loss:
It 18700: Total Loss: 0.19640396535396576, 7.922106742858887	Rec Loss: 0.15679343044757843,	KL Loss:
Done!		







Let's look at some reconstructions and decoded embedding samples for this beta!

```
# [OPTIONAL] visualize VAE reconstructions and samples from the generative model
print("BEST beta = ", beta)
vis_reconstruction(vae_model, randomize=True)
vis_samples(vae_model)
```

```
BEST beta = 0.005
```





## 4. Embedding Space Interpolation [3pt]

As mentioned in the introduction, AEs and VAEs cannot only be used to generate images, but also to learn low-dimensional representations of their inputs. In this final section we will investigate the representations we learned with both models by **interpolating in embedding space** between different images. We will encode two images into their low-dimensional embedding representations, then interpolate these embeddings and reconstruct the result.

```
# Probl-7
nz=32

def get_image_with_label(target_label):
    """Returns a random image from the training set with the requested digit."""
```

```

for img_batch, label_batch in mnist_data_loader:
    for img, label in zip(img_batch, label_batch):
        if label == target_label:
            return img.to(device)

def interpolate_and_visualize(model, tag, start_img, end_img):
    """Encodes images and performs interpolation. Displays decodings."""
    model.eval()    # put model in eval mode to avoid updating batchnorm

    # encode both images into embeddings (use posterior mean for interpolation)
    z_start = model.encoder(start_img[None].reshape(1,784))[..., :nz]
    z_end = model.encoder(end_img[None].reshape(1,784))[..., :nz]

    # compute interpolated latents
    N_INTER_STEPS = 5
    z_inter = [z_start + i/N_INTER_STEPS * (z_end - z_start) for i in
range(N_INTER_STEPS)]

    # decode interpolated embeddings (as a single batch)
    img_inter = model.decoder(torch.cat(z_inter))
    img_inter = img_inter.reshape(-1, 28, 28)

    # reshape result and display interpolation
    vis_imgs = torch.cat([start_img, img_inter, end_img]).reshape(-1,1,28,28)
    fig = plt.figure(figsize = (10, 10))
    ax1 = plt.subplot(111)
    ax1.imshow(torchvision.utils.make_grid(vis_imgs, nrow=N_INTER_STEPS+2,
pad_value=1.))\
                .data.cpu().numpy().transpose(1, 2, 0), cmap='gray')
    plt.title(tag)
    plt.show()

### Interpolation 1
START_LABEL = 2 # ... TODO CHOOSE
END_LABEL = 0 # ... TODO CHOOSE
# sample two training images with given labels
start_img = get_image_with_label(START_LABEL)
end_img = get_image_with_label(END_LABEL)
# visualize interpolations for AE and VAE models
interpolate_and_visualize(ae_model, "Auto-Encoder", start_img, end_img)
interpolate_and_visualize(vae_model, "Variational Auto-Encoder", start_img, end_img)

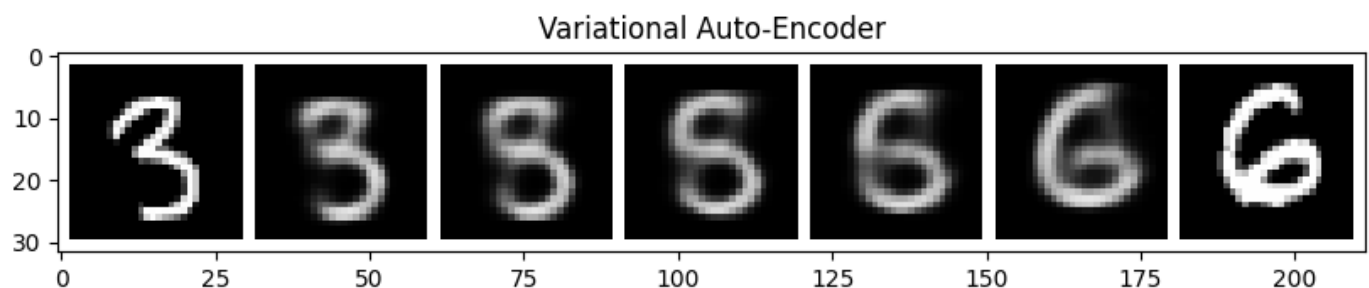



### Interpolation 2
START_LABEL = 3 # ... TODO CHOOSE
END_LABEL = 6 # ... TODO CHOOSE
# sample two training images with given labels
start_img = get_image_with_label(START_LABEL)
end_img = get_image_with_label(END_LABEL)

```

```
# visualize interpolations for AE and VAE models
```

```
interpolate_and_visualize(ae_model, "Auto-Encoder", start_img, end_img)
```

```
interpolate_and_visualize(vae_model, "Variational Auto-Encoder", start_img, end_img)
```



Repeat the experiment for different start / end labels and different samples. Describe your observations.

**Prob1-7 continued: Inline Question: Repeat the interpolation experiment with different start / end labels and multiple samples. Describe your observations! [2 pt]**

1. How do AE and VAE embedding space interpolations differ?
2. How do you expect these differences to affect the usefulness of the learned representation for downstream learning?  
(max 300 words)

**Answer:**

1. From the visualization results, the interpolations in the AE embedding space were generally more linear and less smooth compared to those in the VAE embedding space. Besides, the edges of AE around the digit are not continuous. The VAE embedding space interpolations were more continuous and better followed the underlying data distribution.
2. VAE can get better results than AE in most of the time. In general, VAEs may be more suitable for tasks with complex and difficult-to-model data distributions, while AEs may be more appropriate for simpler, more linear data distributions. Besides, VAE is better suited for learning the underlying data distribution and for generating new data samples due to their ability to sample from the learned latent space distribution during decoding.

## 5. Conditional VAE

Let us now try a Conditional VAE

Now we will try to create a [Conditional VAE](#), where we can condition the encoder and decoder of the VAE on the label `c`.

### Defining the conditional Encoder, Decoder, and VAE models [5 pt]

Prob1-8. We create a separate encoder and decoder class that take in an additional argument `c` in their forward pass, and then build our CVAE model on top of it. Note that the encoder and decoder just need to append `c` to the standard inputs to these modules.

```
def idx2onehot(idx, n):
    """Converts a batch of indices to a one-hot representation."""
    assert torch.max(idx).item() < n
    if idx.dim() == 1:
        idx = idx.unsqueeze(1)
    onehot = torch.zeros(idx.size(0), n).to(idx.device)
    onehot.scatter_(1, idx, 1)

    return onehot

# Let's define encoder and decoder networks

class CVAEEncoder(nn.Module):
    def __init__(self, nz, input_size, conditional, num_labels):
        super().__init__()
```

```

self.input_size = input_size + num_labels if conditional else input_size
self.num_labels = num_labels
self.conditional = conditional

##### TODO
#####
# Create the network architecture using a nn.Sequential module wrapper.
#
# Encoder Architecture:
#
# - input_size -> 256
#
# - ReLU
#
# - 256 -> 64
#
# - ReLU
#
# - 64 -> nz
#
# HINT: Verify the shapes of intermediate layers by running partial networks
#
# (with the next notebook cell) and visualizing the output shapes.
#

#####
self.net = nn.Sequential(
    nn.Linear(self.input_size, 256),
    nn.ReLU(),
    nn.Linear(256, 64),
    nn.ReLU(),
    nn.Linear(64, nz)
)
##### END TODO
#####

def forward(self, x, c=None):
    ##### TODO
    #####
    # If using conditional VAE, concatenate x and a onehot version of c to create
    #
    # the full input. Use function idx2onehot above.
    #

    #####
    if self.conditional:
        c_onehot = idx2onehot(c, self.num_labels) # convert to one-hot
        x = torch.cat([x, c_onehot], dim=-1)

```

```

#####
    return self.net(x)

class CVAEDecoder(nn.Module):
    def __init__(self, nz, output_size, conditional, num_labels):
        super().__init__()
        self.output_size = output_size
        self.conditional = conditional
        self.num_labels = num_labels
        if self.conditional:
            nz = nz + num_labels
        ##### TODO
#####
        # Create the network architecture using a nn.Sequential module wrapper.
#
        # Decoder Architecture (mirrors encoder architecture):
#
        # - nz -> 64
#
        # - ReLU
#
        # - 64 -> 256
#
        # - ReLU
#
        # - 256 -> output_size
#

#####
    self.net = nn.Sequential(
        nn.Linear(nz, 64),
        nn.ReLU(),
        nn.Linear(64, 256),
        nn.ReLU(),
        nn.Linear(256, output_size),
        nn.Sigmoid()
    )
    ##### END TODO
#####

    def forward(self, z, c=None):
        ##### TODO
#####
        # If using conditional VAE, concatenate z and a onehot version of c to create
#
        # the full embedding. Use function idx2onehot above.
#

```

```

#####
    if self.conditional:
        c_onehot = idx2onehot(c, self.num_labels) # convert to one-hot
        z = torch.cat([z, c_onehot], dim=-1)
        ##### END TODO
#####

    return self.net(z).reshape(-1, 1, self.output_size)

class CVAE(nn.Module):
    def __init__(self, nz, beta=1.0, conditional=False, num_labels=0):
        super().__init__()
        if conditional:
            assert num_labels > 0
            self.beta = beta
            self.nz = nz # add
            self.encoder = CVAEEncoder(2*nz, input_size=in_size, conditional=conditional,
num_labels=num_labels)
            self.decoder = CVAEDecoder(nz, output_size=out_size, conditional=conditional,
num_labels=num_labels)

    def forward(self, x, c=None):
        if x.dim() > 2:
            x = x.view(-1, 28*28)

        q = self.encoder(x,c)
        mu, log_sigma = torch.chunk(q, 2, dim=-1)

        # sample latent variable z with reparametrization
        eps = torch.normal(mean=torch.zeros_like(mu), std=torch.ones_like(log_sigma))
        # eps = torch.randn_like(mu) # Alternatively use this
        z = mu + eps * torch.exp(log_sigma)

        # compute reconstruction
        reconstruction = self.decoder(z, c)

        return {'q': q, 'rec': reconstruction, 'c': c}

    def loss(self, x, outputs):
        ##### TODO
#####
        # Implement the loss computation of the VAE.
        #
        # HINT: Your code should implement the following steps:
        #
        #         1. compute the image reconstruction loss, similar to AE loss above
        #

```



```

#          2. compute the KL divergence loss between the inferred posterior
#
#          distribution and a unit Gaussian prior; you can use the provided
#
#          function above for computing the KL divergence between two
Gaussians #
#          parametrized by mean and log_sigma
#
# HINT: Make sure to compute the KL divergence in the correct order since it is
#
#          not symmetric!! ie.  $KL(p, q) \neq KL(q, p)$ 
#

#####
    loss_func = nn.BCELoss(reduction='mean')
    # loss_func = nn.MSELoss()
    # x: [64, 784], y: [64, 1, 784] -> [64, 784]
    rec_loss = loss_func(outputs['rec'].squeeze(1), x)

    # print(outputs['q'].shape)
    mu1 = outputs['q'][:, :self.nz] # [batch_size, nz] = [64, 32]
    log_sigma1 = outputs['q'][:, self.nz:]
    mu2 = torch.zeros_like(mu1).to(device)
    log_sigma2 = torch.zeros_like(log_sigma1).to(device)
    # make it as an scalar instead of a array
    kl_loss = torch.mean(torch.sum(kl_divergence(mu1, log_sigma1, mu2, log_sigma2),
dim=1), dim=0)
    ##### END TODO
#####

    # return weighted objective
    return rec_loss + self.beta * kl_loss, \
        {'rec_loss': rec_loss, 'kl_loss': kl_loss}

def reconstruct(self, x, c=None):
    """Use mean of posterior estimate for visualization reconstruction."""
    ##### TODO
#####
    # This function is used for visualizing reconstructions of our VAE model. To
    #
    # obtain the maximum likelihood estimate we bypass the sampling procedure of
the #
    # inferred latent and instead directly use the mean of the inferred posterior.
    #
    # HINT: encode the input image and then decode the mean of the posterior to
obtain #
    #          the reconstruction.
    #

```

```
#####
    q = self.encoder(x, c)
    mu, log_sigma = torch.chunk(q, 2, dim=-1)
    image = self.decoder(mu, c)
    ##### END TODO
#####
    return image
```

## Setting up the CVAE Training loop

```
learning_rate = 1e-3
nz = 32

##### TODO #####
# Tune the beta parameter to obtain good training results. However, for the #
# initial experiments leave beta = 0 in order to verify our implementation. #
#####
epochs = 5 # works with fewer epochs than AE, VAE. we only test conditional samples.
beta = 0.005
##### END TODO #####

# build CVAE model
conditional = True
cvae_model = CVAE(nz, beta, conditional=conditional, num_labels=10).to(device) #
transfer model to GPU if available
cvae_model = cvae_model.train() # set model in train mode (eg batchnorm params get
updated)

# build optimizer and loss function
##### TODO #####
# Build the optimizer for the cvae_model. We will again use the Adam optimizer with #
# the given learning rate and otherwise default parameters. #
#####
# same as AE
optimizer = torch.optim.Adam(cvae_model.parameters(), lr=learning_rate)
##### END TODO #####

train_it = 0
rec_loss, kl_loss = [], []
print(f"Running {epochs} epochs with {beta=}")
for ep in range(epochs):
    print(f"Run Epoch {ep}")
    ##### TODO #####
    # Implement the main training loop for the model. #
    # If using conditional VAE, remember to pass the conditional variable c to the #
    # forward pass #
    # HINT: Your training loop should sample batches from the data loader, run the #
```

```

#         forward pass of the model, compute the loss, perform the backward pass and #
#         perform one gradient step with the optimizer.                                #
# HINT: Don't forget to erase old gradients before performing the backward pass.      #
# HINT: As before, we will use the loss() function of our model for computing the     #
#         training loss. It outputs the total training loss and a dict containing     #
#         the breakdown of reconstruction and KL loss.                               #
#####
for X_train, y_train in mnist_data_loader:
    X_train, y_train = X_train.to(device), y_train.to(device)
    optimizer.zero_grad()
    X_train = X_train.reshape([batch_size, in_size])
    # falten X_train: [64, 28, 28] -> [64, 784]
    # print(y_train.shape)
    y_pred = cvae_model(X_train, y_train)
    # print(y_pred['rec'].shape, X_train.shape)
    total_loss, losses = cvae_model.loss(X_train, y_pred)
    # losses['rec_loss'] = losses['rec_loss'].detach().cpu()
    # losses['kl_loss'] = losses['kl_loss'].detach().cpu()
    # print(total_loss.shape)

    total_loss.backward()
    optimizer.step()

    rec_loss.append(losses['rec_loss'])
    kl_loss.append(losses['kl_loss'])
    if train_it % 100 == 0:
        print("It {}: Total Loss: {}, \t Rec Loss: {},\t KL Loss: {}".format(
            train_it, total_loss, losses['rec_loss'], losses['kl_loss']))
    train_it += 1
##### END TODO #####

print("Done!")

rec_loss_plotdata = [foo.detach().cpu() for foo in rec_loss]
kl_loss_plotdata = [foo.detach().cpu() for foo in kl_loss]

# log the loss training curves
fig = plt.figure(figsize = (10, 5))
ax1 = plt.subplot(121)
ax1.plot(rec_loss_plotdata)
ax1.title.set_text("Reconstruction Loss")
ax2 = plt.subplot(122)
ax2.plot(kl_loss_plotdata)
ax2.title.set_text("KL Loss")
plt.show()

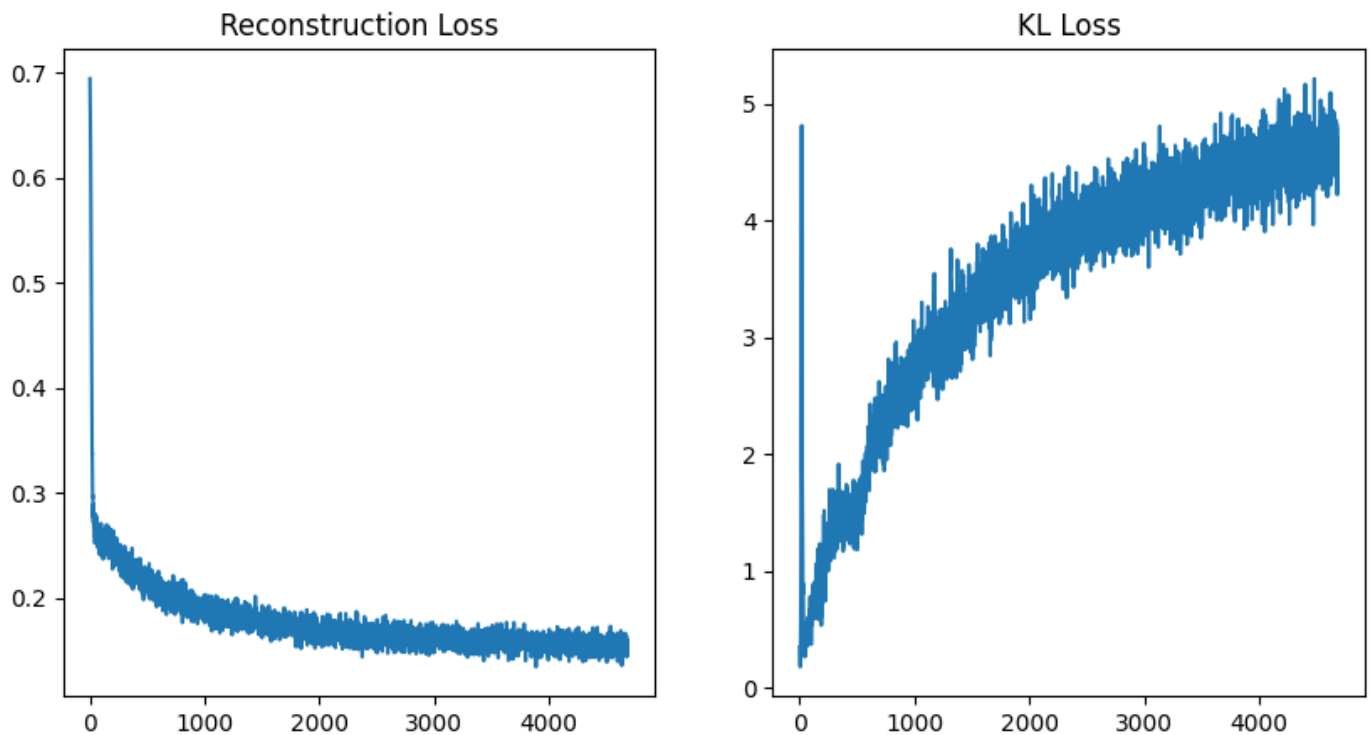
```

Running 5 epochs with beta=0.005  
Run Epoch 0

It 0: Total Loss: 0.6961221098899841, 0.3461534380912781	Rec Loss: 0.6943913698196411,	KL Loss:
It 100: Total Loss: 0.26270538568496704, 0.4554291367530823	Rec Loss: 0.260428249835968,	KL Loss:
It 200: Total Loss: 0.23876087367534637, 0.797675371170044	Rec Loss: 0.23477250337600708,	KL Loss:
It 300: Total Loss: 0.2432468831539154, 1.4952359199523926	Rec Loss: 0.23577070236206055,	KL Loss:
It 400: Total Loss: 0.22452087700366974, 1.4124525785446167	Rec Loss: 0.21745862066745758,	KL Loss:
It 500: Total Loss: 0.22840912640094757, 1.1791630983352661	Rec Loss: 0.2225133180618286,	KL Loss:
It 600: Total Loss: 0.21754541993141174, 1.9364993572235107	Rec Loss: 0.20786292850971222,	KL Loss:
It 700: Total Loss: 0.20954541862010956, 2.0965616703033447	Rec Loss: 0.19906261563301086,	KL Loss:
It 800: Total Loss: 0.2016667127609253, 2.1582891941070557	Rec Loss: 0.19087526202201843,	KL Loss:
It 900: Total Loss: 0.20380091667175293, 2.427638530731201	Rec Loss: 0.1916627287864685,	KL Loss:
Run Epoch 1		
It 1000: Total Loss: 0.1933600753545761, 2.5831735134124756	Rec Loss: 0.18044421076774597,	KL Loss:
It 1100: Total Loss: 0.1996302604675293, 3.0014185905456543	Rec Loss: 0.18462316691875458,	KL Loss:
It 1200: Total Loss: 0.190623477101326, 2.4660449028015137	Rec Loss: 0.17829325795173645,	KL Loss:
It 1300: Total Loss: 0.1867140382528305, 2.9099395275115967	Rec Loss: 0.17216433584690094,	KL Loss:
It 1400: Total Loss: 0.18920038640499115, 2.737218141555786	Rec Loss: 0.17551429569721222,	KL Loss:
It 1500: Total Loss: 0.17712968587875366, 3.08811092376709	Rec Loss: 0.1616891324520111,	KL Loss:
It 1600: Total Loss: 0.19290322065353394, 3.269814968109131	Rec Loss: 0.1765541434288025,	KL Loss:
It 1700: Total Loss: 0.19574669003486633, 3.866971254348755	Rec Loss: 0.17641183733940125,	KL Loss:
It 1800: Total Loss: 0.1918659508228302, 3.4042232036590576	Rec Loss: 0.17484483122825623,	KL Loss:
Run Epoch 2		
It 1900: Total Loss: 0.1924508512020111, 3.551847457885742	Rec Loss: 0.17469161748886108,	KL Loss:
It 2000: Total Loss: 0.18140479922294617, 3.297797203063965	Rec Loss: 0.16491581499576569,	KL Loss:
It 2100: Total Loss: 0.18916717171669006, 3.715733289718628	Rec Loss: 0.17058850824832916,	KL Loss:
It 2200: Total Loss: 0.19339223206043243, 4.19720458984375	Rec Loss: 0.17240621149539948,	KL Loss:

It 2300: Total Loss: 0.18865463137626648, 3.9523258209228516	Rec Loss: 0.1688929945230484,	KL Loss:
It 2400: Total Loss: 0.1980040967464447, 4.018292427062988	Rec Loss: 0.177912637591362,	KL Loss:
It 2500: Total Loss: 0.19335682690143585, 3.734665870666504	Rec Loss: 0.17468349635601044,	KL Loss:
It 2600: Total Loss: 0.18068158626556396, 3.8151142597198486	Rec Loss: 0.16160601377487183,	KL Loss:
It 2700: Total Loss: 0.18145106732845306, 4.070026397705078	Rec Loss: 0.16110093891620636,	KL Loss:
It 2800: Total Loss: 0.18490146100521088, 4.102055549621582	Rec Loss: 0.1643911898136139,	KL Loss:
Run Epoch 3		
It 2900: Total Loss: 0.18923825025558472, 4.087955474853516	Rec Loss: 0.16879847645759583,	KL Loss:
It 3000: Total Loss: 0.17653420567512512, 3.9337525367736816	Rec Loss: 0.1568654477596283,	KL Loss:
It 3100: Total Loss: 0.17404583096504211, 4.123428821563721	Rec Loss: 0.15342868864536285,	KL Loss:
It 3200: Total Loss: 0.17988714575767517, 4.348489761352539	Rec Loss: 0.15814469754695892,	KL Loss:
It 3300: Total Loss: 0.17844687402248383, 4.142595291137695	Rec Loss: 0.15773390233516693,	KL Loss:
It 3400: Total Loss: 0.17906738817691803, 4.371393203735352	Rec Loss: 0.15721042454242706,	KL Loss:
It 3500: Total Loss: 0.18766286969184875, 4.235945224761963	Rec Loss: 0.16648314893245697,	KL Loss:
It 3600: Total Loss: 0.1857607513666153, 4.351469039916992	Rec Loss: 0.1640034019947052,	KL Loss:
It 3700: Total Loss: 0.18224720656871796, 4.154415130615234	Rec Loss: 0.16147513687610626,	KL Loss:
Run Epoch 4		
It 3800: Total Loss: 0.186319962143898, 4.644988059997559	Rec Loss: 0.16309502720832825,	KL Loss:
It 3900: Total Loss: 0.17574700713157654, 4.3214921951293945	Rec Loss: 0.15413954854011536,	KL Loss:
It 4000: Total Loss: 0.19297140836715698, 4.616983890533447	Rec Loss: 0.1698864847421646,	KL Loss:
It 4100: Total Loss: 0.17652323842048645, 4.431282043457031	Rec Loss: 0.15436682105064392,	KL Loss:
It 4200: Total Loss: 0.18403884768486023, 4.662446022033691	Rec Loss: 0.1607266217470169,	KL Loss:
It 4300: Total Loss: 0.1708582639694214, 4.541799545288086	Rec Loss: 0.1481492668390274,	KL Loss:
It 4400: Total Loss: 0.18357664346694946, 4.892752647399902	Rec Loss: 0.15911288559436798,	KL Loss:
It 4500: Total Loss: 0.17335020005702972, 4.349187850952148	Rec Loss: 0.15160426497459412,	KL Loss:

It 4600: Total Loss: 0.17443019151687622, Rec Loss: 0.1515093594789505, KL Loss: 4.58416748046875  
Done!




## Verifying conditional samples from CVAE [6 pt]

Now let us generate samples from the trained model, conditioned on all the labels.

```
# Probl-9
if conditional:
    c = torch.arange(0, 10).long().unsqueeze(1).to(device)
    z = torch.randn([10, nz]).to(device)
    x = cvae_model.decoder(z, c=c)
else:
    z = torch.randn([10, nz]).to(device)
    x = cvae_model.decoder(z)

plt.figure()
plt.figure(figsize=(5, 10))
for p in range(10):
    plt.subplot(5, 2, p+1)
    if conditional:
        plt.text(
            0, 0, "c={:d}".format(c[p].item()), color='black',
            backgroundColor='white', fontsize=8)
    plt.imshow(x[p].view(28, 28).cpu().data.numpy(), cmap='gray')
    plt.axis('off')
```

<Figure size 640x480 with 0 Axes>



## Submission Instructions

You need to submit this jupyter notebook and a PDF. See Piazza for detailed submission instructions.